

# Package ‘asa’

June 22, 2026

**Title** AI Search Agent for Large-Scale Research Automation

**Version** 0.1.0

**Author** Connor Jerzak [aut, cre] (<<https://orcid.org/0000-0003-1914-8905>>)

**Maintainer** Connor Jerzak <connor.jerzak@gmail.com>

**Description** Provides an LLM-powered research agent for performing AI search tasks at large scales. Uses a ReAct (Reasoning + Acting) agent pattern with web search capabilities via DuckDuckGo and Wikipedia. Implements DeepAgent-style memory folding for context management. Agent runtimes include default 'OpenCode' and a built-in 'LangGraph' runtime. The package supports multiple LLM backends including 'OpenAI', 'Groq', 'xAI', 'Gemini', 'Anthropic' (Claude), 'AWS Bedrock', 'OpenRouter', and 'Exo'.

**URL** <https://github.com/cjerzak/asa-software>

**BugReports** <https://github.com/cjerzak/asa-software/issues>

**Depends** R (>= 4.0.0)

**License** MIT + file LICENSE

**Encoding** UTF-8

**LazyData** false

**Imports** reticulate (>= 1.28),  
jsonlite,  
rlang,  
digest,  
yaml,  
processx

**Suggests** testthat (>= 3.0.0),  
knitr,  
rmarkdown,  
future,  
future.apply,  
withr,  
DBI,  
RSQLite

**VignetteBuilder** knitr

**Config/testthat/edition** 3

**SystemRequirements** Python (>= 3.11), Conda, OpenCode CLI (default agent runtime; optional when using another agent\_backend), Tor (optional, for anonymous searching)

**Config/roxygen2/version 8.0.0**

## Contents

|                                               |    |
|-----------------------------------------------|----|
| as.data.frame.asa_audit_result . . . . .      | 3  |
| as.data.frame.asa_enumerate_result . . . . .  | 4  |
| as.data.frame.asa_evaluation_result . . . . . | 4  |
| as.data.frame.asa_result . . . . .            | 5  |
| asa_agent . . . . .                           | 5  |
| asa_audit . . . . .                           | 6  |
| asa_audit_result . . . . .                    | 8  |
| asa_config . . . . .                          | 9  |
| asa_enumerate . . . . .                       | 11 |
| asa_enumerate_result . . . . .                | 15 |
| asa_evaluation_result . . . . .               | 16 |
| asa_response . . . . .                        | 17 |
| asa_result . . . . .                          | 18 |
| benchmark_from_df . . . . .                   | 19 |
| benchmark_from_yaml . . . . .                 | 20 |
| build_backend . . . . .                       | 21 |
| build_prompt . . . . .                        | 22 |
| check_backend . . . . .                       | 23 |
| configure_search . . . . .                    | 24 |
| configure_search_logging . . . . .            | 25 |
| configure_temporal . . . . .                  | 26 |
| configure_tor_registry . . . . .              | 27 |
| create_benchmark . . . . .                    | 28 |
| evaluate_agent . . . . .                      | 29 |
| extract_agent_results . . . . .               | 29 |
| extract_search_snippets . . . . .             | 30 |
| extract_search_tiers . . . . .                | 31 |
| extract_urls . . . . .                        | 32 |
| extract_wikipedia_content . . . . .           | 32 |
| get_agent . . . . .                           | 33 |
| get_tor_ip . . . . .                          | 33 |
| initialize_agent . . . . .                    | 34 |
| is_tor_running . . . . .                      | 37 |
| list_benchmarks . . . . .                     | 38 |
| load_benchmark . . . . .                      | 38 |
| print.asa_agent . . . . .                     | 39 |
| print.asa_audit_result . . . . .              | 39 |
| print.asa_benchmark . . . . .                 | 40 |
| print.asa_config . . . . .                    | 40 |
| print.asa_enumerate_result . . . . .          | 41 |
| print.asa_evaluation_result . . . . .         | 41 |
| print.asa_response . . . . .                  | 42 |
| print.asa_result . . . . .                    | 42 |
| print.asa_search . . . . .                    | 43 |
| print.asa_temporal . . . . .                  | 43 |
| print.asa_tor . . . . .                       | 44 |
| process_outputs . . . . .                     | 44 |
| reset_agent . . . . .                         | 45 |



---

```
as.data.frame.asa_enumerate_result
```

*Convert asa\_enumerate\_result to Data Frame*

---

### Description

Convert asa\_enumerate\_result to Data Frame

### Usage

```
## S3 method for class 'asa_enumerate_result'  
as.data.frame(x, ...)
```

### Arguments

|     |                                |
|-----|--------------------------------|
| x   | An asa_enumerate_result object |
| ... | Additional arguments (ignored) |

### Value

The data data.frame from the result

---

```
as.data.frame.asa_evaluation_result
```

*Convert asa\_evaluation\_result to Data Frame*

---

### Description

Convert asa\_evaluation\_result to Data Frame

### Usage

```
## S3 method for class 'asa_evaluation_result'  
as.data.frame(x, ...)
```

### Arguments

|     |                                 |
|-----|---------------------------------|
| x   | An asa_evaluation_result object |
| ... | Additional arguments (ignored)  |

### Value

The per-task evaluation records as a data.frame

---

|                          |                                         |
|--------------------------|-----------------------------------------|
| as.data.frame.asa_result | <i>Convert asa_result to Data Frame</i> |
|--------------------------|-----------------------------------------|

---

**Description**

Convert asa\_result to Data Frame

**Usage**

```
## S3 method for class 'asa_result'  
as.data.frame(x, ...)
```

**Arguments**

- x                    An asa\_result object
- ...                  Additional arguments (ignored)

**Value**

A single-row data frame

---

|           |                                          |
|-----------|------------------------------------------|
| asa_agent | <i>Constructor for asa_agent Objects</i> |
|-----------|------------------------------------------|

---

**Description**

Creates an S3 object representing an initialized ASA search agent.

**Usage**

```
asa_agent(python_agent, backend, model, config, llm = NULL, tools = NULL)
```

**Arguments**

- python\_agent        The underlying Python agent object
- backend             LLM backend name (e.g., "openai", "groq")
- model               Model identifier
- config              Agent configuration list
- llm                  Optional LLM object used by LangGraph
- tools                Optional list of tools associated with the agent

**Value**

An object of class asa\_agent

asa\_audit

*Audit Enumeration Results for Completeness and Quality***Description**

Validates enumeration results for completeness, consistency, and data quality using either Claude Code (CLI) or a LangGraph-based audit pipeline.

**Usage**

```
asa_audit(
  result,
  query = NULL,
  known_universe = NULL,
  checks = c("completeness", "consistency", "gaps", "anomalies"),
  backend = c("claude_code", "langgraph"),
  claude_model = "claude-sonnet-4-20250514",
  llm_model = "gpt-4.1-mini",
  interactive = FALSE,
  confidence_threshold = 0.8,
  timeout = 120,
  verbose = TRUE,
  agent = NULL
)
```

**Arguments**

|                      |                                                                                                                                 |
|----------------------|---------------------------------------------------------------------------------------------------------------------------------|
| result               | An <code>asa_enumerate_result</code> object or a <code>data.frame</code> to audit                                               |
| query                | The original enumeration query (inferred from result if NULL)                                                                   |
| known_universe       | Optional vector of expected items for completeness check                                                                        |
| checks               | Character vector of checks to perform. Options: "completeness", "consistency", "gaps", "anomalies". Default runs all checks.    |
| backend              | Backend to use for auditing: "claude_code" (CLI) or "langgraph"                                                                 |
| claude_model         | Model to use with Claude Code backend                                                                                           |
| llm_model            | Model to use with LangGraph backend                                                                                             |
| interactive          | If TRUE and using <code>claude_code</code> backend, spawn an interactive Claude Code session instead of programmatic invocation |
| confidence_threshold | Flag items with confidence below this threshold                                                                                 |
| timeout              | Timeout in seconds for the audit operation                                                                                      |
| verbose              | Print progress messages                                                                                                         |
| agent                | Existing <code>asa_agent</code> for LangGraph backend (optional)                                                                |

## Details

The audit function adds three columns to the data:

- `_audit_flag`: "ok", "warning", or "suspect"
- `_audit_notes`: Explanation of any issues
- `_confidence_adjusted`: Revised confidence after audit

### ## Audit Checks

**completeness**: Checks for missing items by comparing against `known_universe` (if provided) or using domain knowledge.

**consistency**: Validates data types, patterns, and value ranges.

**gaps**: Identifies systematic patterns of missing data (geographic, temporal, categorical gaps).

**anomalies**: Detects duplicates, outliers, and suspicious patterns.

## Value

An `asa_audit_result` object containing:

|                                 |                                                                                                 |
|---------------------------------|-------------------------------------------------------------------------------------------------|
| <code>data</code>               | Original data with audit columns added ( <code>_audit_flag</code> , <code>_audit_notes</code> ) |
| <code>audit_summary</code>      | High-level summary of findings                                                                  |
| <code>issues</code>             | List of identified issues with severity and descriptions                                        |
| <code>recommendations</code>    | Suggested remediation queries                                                                   |
| <code>completeness_score</code> | 0-1 score for data completeness                                                                 |
| <code>consistency_score</code>  | 0-1 score for data consistency                                                                  |

## Examples

```
## Not run:
# Audit enumeration results with Claude Code
senators <- asa_enumerate(
  query = "Find all current US senators",
  schema = c(name = "character", state = "character", party = "character")
)
audit <- asa_audit(senators, backend = "claude_code")
print(audit)

# Audit with known universe for precise completeness check
audit <- asa_audit(senators, known_universe = state.abb)

# Interactive mode for complex audits
asa_audit(senators, backend = "claude_code", interactive = TRUE)

# Use LangGraph backend
audit <- asa_audit(senators, backend = "langgraph", agent = agent)

## End(Not run)
```

---

|                  |                                                 |
|------------------|-------------------------------------------------|
| asa_audit_result | <i>Constructor for asa_audit_result Objects</i> |
|------------------|-------------------------------------------------|

---

## Description

Creates an S3 object representing the result of a data quality audit.

## Usage

```
asa_audit_result(  
  data,  
  audit_summary,  
  issues,  
  recommendations,  
  completeness_score,  
  consistency_score,  
  backend_used,  
  elapsed_time,  
  query = NULL,  
  checks = NULL  
)
```

## Arguments

|                    |                                                                              |
|--------------------|------------------------------------------------------------------------------|
| data               | data.frame with original data plus audit columns (_audit_flag, _audit_notes) |
| audit_summary      | Character string with high-level findings                                    |
| issues             | List of identified issues with severity and descriptions                     |
| recommendations    | Character vector of suggested remediation queries                            |
| completeness_score | Numeric 0-1 score for data completeness                                      |
| consistency_score  | Numeric 0-1 score for data consistency                                       |
| backend_used       | Which backend performed the audit ("claude_code" or "langgraph")             |
| elapsed_time       | Execution time in seconds                                                    |
| query              | The original query (if available)                                            |
| checks             | Character vector of checks that were performed                               |

## Value

An object of class `asa_audit_result`



asa\_config

*Create ASA Configuration Object***Description**

Creates a configuration object that encapsulates all settings for ASA tasks. This provides a unified way to configure backend, model, search, temporal, and resource settings in a single object.

**Usage**

```
asa_config(
  agent_backend = NULL,
  backend = NULL,
  model = NULL,
  conda_env = NULL,
  proxy = NA,
  use_browser = NULL,
  workers = NULL,
  timeout = NULL,
  rate_limit = NULL,
  memory_folding = NULL,
  memory_threshold = NULL,
  memory_keep_recent = NULL,
  use_observational_memory = NULL,
  om_observation_token_budget = NULL,
  om_reflection_token_budget = NULL,
  om_buffer_tokens = NULL,
  om_buffer_activation = NULL,
  om_block_after = NULL,
  om_async_prebuffer = NULL,
  om_cross_thread_memory = NULL,
  temporal = NULL,
  search = NULL,
  tor = NULL,
  recursion_limit = NULL
)
```

**Arguments**

|               |                                                                                                                                                                                                                                                                          |
|---------------|--------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| agent_backend | Agent runtime backend: "opencode" (default OpenCode CLI backend with ASA-managed provider routing and search MCP tools), "agent" (built-in ASA Lang-Graph pipeline), or "free-code" (headless free-code backend with ASA-managed provider routing and search MCP tools). |
| backend       | LLM backend: "openai", "azure-openai", "groq", "xai", "gemini", "exo", "ollama", "openrouter", "anthropic", or "bedrock". For Azure OpenAI, model is the Azure deployment name. Ollama also supports shorthand like "ollama-1fm2:24b-a2b" when model is omitted.         |
| model         | Model identifier (e.g., "gpt-4.1-mini")                                                                                                                                                                                                                                  |
| conda_env     | Conda environment name. Defaults to the package option <code>asa.default_conda_env</code> (or "asa_env" if unset).                                                                                                                                                       |

|                             |                                                                                                                                                                      |
|-----------------------------|----------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| proxy                       | Proxy URL (e.g., Tor SOCKS5). Use NA (default) to auto-detect from environment variables (ASA_PROXY, HTTP_PROXY, HTTPS_PROXY); use NULL to disable proxying.         |
| use_browser                 | Enable Selenium browser tier for DuckDuckGo search.                                                                                                                  |
| workers                     | Number of parallel workers for batch operations                                                                                                                      |
| timeout                     | Request timeout in seconds                                                                                                                                           |
| rate_limit                  | Requests per second                                                                                                                                                  |
| memory_folding              | Enable DeepAgent-style memory folding                                                                                                                                |
| memory_threshold            | Messages before folding triggers                                                                                                                                     |
| memory_keep_recent          | Exchanges to preserve after folding. An exchange is a user turn plus the assistant response, including any tool calls and tool outputs.                              |
| use_observational_memory    | Enable observational memory pipeline.                                                                                                                                |
| om_observation_token_budget | Token budget for stored observation messages.                                                                                                                        |
| om_reflection_token_budget  | Token budget for reflection summaries.                                                                                                                               |
| om_buffer_tokens            | Token budget for buffered context before reflection.                                                                                                                 |
| om_buffer_activation        | Fraction of buffer usage that triggers reflection.                                                                                                                   |
| om_block_after              | Fraction of context usage where buffering blocks generation.                                                                                                         |
| om_async_prebuffer          | Whether to pre-buffer observations asynchronously.                                                                                                                   |
| om_cross_thread_memory      | Whether observational memory is shared across threads.                                                                                                               |
| temporal                    | Temporal filtering options (use temporal_options())                                                                                                                  |
| search                      | Search configuration (use search_options())                                                                                                                          |
| tor                         | Tor registry options (use tor_options())                                                                                                                             |
| recursion_limit             | Optional default maximum number of agent steps. When NULL, mode-specific defaults are used at runtime unless overridden in <a href="#">run_task/run_task_batch</a> . |

## Details

The configuration object can be passed to `run_task()`, `run_task_batch()`, `asa_enumerate()`, and other functions to provide consistent settings across operations.

## Value

An object of class `asa_config`

## See Also

[temporal\\_options](#), [search\\_options](#)

## Examples

```
## Not run:
# Create configuration
config <- asa_config(
  backend = "openai",
  model = "gpt-4.1-mini",
  workers = 4,
  temporal = temporal_options(time_filter = "y")
)

# Use with run_task
result <- run_task("Find recent AI breakthroughs", config = config)

## End(Not run)
```

asa\_enumerate

*Multi-Agent Research for Open-Ended Queries*

## Description

Performs intelligent open-ended research tasks using multi-agent orchestration. Decomposes complex queries into sub-tasks, executes parallel searches, and aggregates results into structured output (data.frame, CSV, or JSON).

## Usage

```
asa_enumerate(
  query = NULL,
  schema = NULL,
  output = c("data.frame", "csv", "json"),
  config = NULL,
  workers = NULL,
  max_rounds = NULL,
  budget = list(queries = 50L, tokens = 200000L, time_sec = 300L),
  stop_policy = list(target_items = NULL, plateau_rounds = 2L, novelty_min = 0.05,
    novelty_window = 20L),
  sources = list(web = TRUE, wikipedia = TRUE, wikidata = TRUE),
  allow_read_webpages = NULL,
  webpage_relevance_mode = NULL,
  webpage_embedding_provider = NULL,
  webpage_embedding_model = NULL,
  temporal = NULL,
  pagination = TRUE,
  progress = TRUE,
  include_provenance = FALSE,
  checkpoint = TRUE,
  checkpoint_dir = tempdir(),
  resume_from = NULL,
  agent = NULL,
  backend = NULL,
  model = NULL,
```

```

    conda_env = NULL,
    verbose = TRUE
)

```

## Arguments

|                            |                                                                                                                                                                                                                                                                                                                                                                                    |
|----------------------------|------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| query                      | Character string describing the research goal. Required for fresh runs; optional when resume_from points to an existing checkpoint. Examples: "Find all current US senators with their state, party, and term end date"                                                                                                                                                            |
| schema                     | Named character vector defining the output schema. Names are column names, values are R types ("character", "numeric", "logical"). Use NULL or "auto" for LLM-proposed schema.                                                                                                                                                                                                     |
| output                     | Output format: "data.frame" (default), "csv", or "json".                                                                                                                                                                                                                                                                                                                           |
| config                     | Optional asa_config object to supply defaults for agent initialization and per-run search/temporal settings. Explicit arguments to asa_enumerate() take precedence over config.                                                                                                                                                                                                    |
| workers                    | Number of parallel search workers. Defaults to value from config\$workers when provided, otherwise ASA_DEFAULT_WORKERS (typically 4).                                                                                                                                                                                                                                              |
| max_rounds                 | Maximum research iterations. Defaults to value from ASA_DEFAULT_MAX_ROUNDS (typically 8).                                                                                                                                                                                                                                                                                          |
| budget                     | Named list with resource limits: <ul style="list-style-type: none"> <li>queries: Maximum search queries (default: 50)</li> <li>tokens: Maximum LLM tokens (default: 200000)</li> <li>time_sec: Maximum execution time in seconds (default: 300)</li> </ul>                                                                                                                         |
| stop_policy                | Named list with stopping criteria: <ul style="list-style-type: none"> <li>target_items: Stop when this many items found (NULL = unknown)</li> <li>plateau_rounds: Stop after N rounds with no new items (default: 2)</li> <li>novelty_min: Minimum new items ratio per round (default: 0.05)</li> <li>novelty_window: Window size for novelty calculation (default: 20)</li> </ul> |
| sources                    | Named list controlling which sources to use: <ul style="list-style-type: none"> <li>web: Use DuckDuckGo web search (default: TRUE)</li> <li>wikipedia: Use Wikipedia (default: TRUE)</li> <li>wikidata: Use Wikidata SPARQL for authoritative enumerations (default: TRUE)</li> </ul>                                                                                              |
| allow_read_webpages        | If TRUE, the agent may open and read full webpages (in addition to search snippets) when it helps extraction. Use NULL (default) to inherit from config\$search (when provided); otherwise defaults to FALSE for safety and to avoid large context usage.                                                                                                                          |
| webpage_relevance_mode     | Relevance selection for opened webpages. One of: "auto" (default), "lexical", "embeddings". When "embeddings" or "auto" with an available provider, the tool uses vector similarity to pick the most relevant excerpts; otherwise it falls back to lexical overlap.                                                                                                                |
| webpage_embedding_provider | Embedding provider to use for relevance. One of: "auto" (default), "openai", "sentence_transformers".                                                                                                                                                                                                                                                                              |

|                         |                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                            |
|-------------------------|--------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| webpage_embedding_model | Embedding model identifier. For OpenAI, defaults to "text-embedding-3-small". For sentence-transformers, use a local model name (e.g., "all-MiniLM-L6-v2").                                                                                                                                                                                                                                                                                                                                                                                |
| temporal                | Named list for temporal filtering: <ul style="list-style-type: none"> <li>• after: ISO 8601 date string (e.g., "2020-01-01") - results after this date</li> <li>• before: ISO 8601 date string (e.g., "2024-01-01") - results before this date</li> <li>• time_filter: DuckDuckGo time filter ("d", "w", "m", "y") for day/week/month/year</li> <li>• strictness: "best_effort" (default) or "strict" (verifies dates via metadata)</li> <li>• use_wayback: Use Wayback Machine for strict pre-date guarantees (default: FALSE)</li> </ul> |
| pagination              | Enable pagination for large result sets (default: TRUE).                                                                                                                                                                                                                                                                                                                                                                                                                                                                                   |
| progress                | Show progress bar and status updates (default: TRUE).                                                                                                                                                                                                                                                                                                                                                                                                                                                                                      |
| include_provenance      | Include source URLs and confidence per row (default: FALSE).                                                                                                                                                                                                                                                                                                                                                                                                                                                                               |
| checkpoint              | Enable auto-save after planning and each completed round (default: TRUE).                                                                                                                                                                                                                                                                                                                                                                                                                                                                  |
| checkpoint_dir          | Directory for checkpoint files (default: tempdir()).                                                                                                                                                                                                                                                                                                                                                                                                                                                                                       |
| resume_from             | Path to checkpoint file to resume from (default: NULL). For v2 checkpoints this continues execution from the last saved round state. Legacy v1 checkpoints return the saved result without continuation.                                                                                                                                                                                                                                                                                                                                   |
| agent                   | An initialized asa_agent object. If NULL, uses the current agent or creates a new one with specified backend/model.                                                                                                                                                                                                                                                                                                                                                                                                                        |
| backend                 | LLM backend if creating new agent: "openai", "groq", "xai", "gemini", "exo", "ollama", "openrouter", "anthropic", or "bedrock".                                                                                                                                                                                                                                                                                                                                                                                                            |
| model                   | Model identifier if creating new agent.                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                    |
| conda_env               | Conda environment name. Defaults to the package option asa.default_conda_env (or "asa_env" if unset).                                                                                                                                                                                                                                                                                                                                                                                                                                      |
| verbose                 | Print status messages (default: TRUE).                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                     |

## Details

The function uses an iterative graph architecture:

1. **Planner:** Decomposes the query and proposes a search plan.
2. **Searcher:** Queries Wikidata (when applicable) and falls back to web/Wikipedia.
3. **Deduper:** Removes duplicates using hashing + fuzzy matching.
4. **Stopper:** Evaluates stopping criteria (novelty, budgets, saturation).

Parallelism is limited and backend-dependent. For example, strict temporal filtering may verify publication dates in parallel up to workers.

For known entity types (US senators, countries, Fortune 500), Wikidata provides authoritative enumerations with complete, verified data.

## Value

An object of class `asa_enumerate_result` containing:

- data: data.frame with results matching the schema
- status: "complete", "partial", or "failed"

- stop\_reason: Why the search stopped
- metrics: List with rounds, queries\_used, novelty\_curve, coverage
- provenance: If include\_provenance=TRUE, source info per row
- checkpoint\_file: Path to checkpoint if saved

### Checkpointing

With checkpoint=TRUE, state is saved after planning and after each completed round. If interrupted, use resume\_from to continue from the last checkpoint:

```
result <- asa_enumerate(resume_from = "/path/to/checkpoint.rds")
```

### Schema

The schema defines expected output columns:

```
schema = c(name = "character", state = "character", party = "character")
```

With schema = "auto", the planner agent proposes a schema based on the query.

### See Also

[run\\_task](#), [initialize\\_agent](#)

### Examples

```
## Not run:
# Find all US senators
senators <- asa_enumerate(
  query = "Find all current US senators with state, party, and term end date",
  schema = c(name = "character", state = "character",
              party = "character", term_end = "character"),
  stop_policy = list(target_items = 100),
  include_provenance = TRUE
)
head(senators$data)

# Find countries with auto schema
countries <- asa_enumerate(
  query = "Find all countries with their capitals and populations",
  schema = "auto",
  output = "csv"
)

# Resume from checkpoint
result <- asa_enumerate(
  query = "Find Fortune 500 CEOs",
  resume_from = "/tmp/asa_enumerate_abc123.rds"
)

# Temporal filtering: results from specific date range
companies_2020s <- asa_enumerate(
  query = "Find tech companies founded recently",
  temporal = list(
    after = "2020-01-01",
```

```

        before = "2024-01-01",
        strictness = "best_effort"
    )
)

# Temporal filtering: past year with DuckDuckGo time filter
recent_news <- asa_enumerate(
  query = "Find AI research breakthroughs",
  temporal = list(
    time_filter = "y" # past year
  )
)

# Strict temporal filtering with Wayback Machine
historical <- asa_enumerate(
  query = "Find Fortune 500 companies",
  temporal = list(
    before = "2015-01-01",
    strictness = "strict",
    use_wayback = TRUE
  )
)

## End(Not run)

```

---

asa\_enumerate\_result    *Constructor for asa\_enumerate\_result Objects*

---

## Description

Creates an S3 object representing the result of an enumeration task.

## Usage

```

asa_enumerate_result(
  data,
  status,
  stop_reason,
  metrics,
  provenance = NULL,
  plan = NULL,
  checkpoint_file = NULL,
  query = NULL,
  schema = NULL
)

```

## Arguments

|             |                                                                         |
|-------------|-------------------------------------------------------------------------|
| data        | data.frame containing the enumeration results                           |
| status      | Result status: "complete", "partial", or "failed"                       |
| stop_reason | Why the enumeration stopped (e.g., "target_reached", "novelty_plateau") |

|                 |                                                          |
|-----------------|----------------------------------------------------------|
| metrics         | List with execution metrics (rounds, queries_used, etc.) |
| provenance      | Optional data.frame with source information per row      |
| plan            | The enumeration plan from the planner agent              |
| checkpoint_file | Path to saved checkpoint file                            |
| query           | The original enumeration query                           |
| schema          | The schema used for extraction                           |

**Value**

An object of class `asa_enumerate_result`

---

`asa_evaluation_result` *Constructor for asa\_evaluation\_result Objects*

---

**Description**

Creates an S3 object representing the result of an agent evaluation run.

**Usage**

```
asa_evaluation_result(
  benchmark,
  task_results,
  run_summaries,
  metrics,
  agent_spec,
  config_snapshot,
  outputs = NULL,
  elapsed_time = NA_real_
)
```

**Arguments**

|                 |                                               |
|-----------------|-----------------------------------------------|
| benchmark       | Benchmark object used for the evaluation      |
| task_results    | data.frame with one row per task execution    |
| run_summaries   | data.frame with one row per evaluation run    |
| metrics         | Aggregate evaluation metrics                  |
| agent_spec      | Named list describing the evaluated agent     |
| config_snapshot | Configuration snapshot used during evaluation |
| outputs         | Optional named list of retained raw outputs   |
| elapsed_time    | Total elapsed evaluation time in minutes      |

**Value**

An object of class `asa_evaluation_result`



---

|              |                                             |
|--------------|---------------------------------------------|
| asa_response | <i>Constructor for asa_response Objects</i> |
|--------------|---------------------------------------------|

---

## Description

Creates an S3 object representing an agent response.

## Usage

```
asa_response(
  message,
  status_code,
  raw_response,
  trace,
  elapsed_time,
  prompt,
  trace_json = "",
  fold_stats = list(),
  thread_id = NULL,
  stop_reason = NULL,
  budget_state = list(),
  field_status = list(),
  diagnostics = list(),
  json_repair = list(),
  final_payload = NULL,
  terminal_valid = FALSE,
  terminal_payload_hash = NULL,
  payload_integrity = list(),
  completion_gate = list(),
  tokens_used = NA_integer_,
  input_tokens = NA_integer_,
  output_tokens = NA_integer_,
  token_trace = list()
)
```

## Arguments

|              |                                                                                                                                                                                                                                                                                                                                                                                                                       |
|--------------|-----------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| message      | The final response text                                                                                                                                                                                                                                                                                                                                                                                               |
| status_code  | Status code (200 = success, 100 = error)                                                                                                                                                                                                                                                                                                                                                                              |
| raw_response | The full Python response object                                                                                                                                                                                                                                                                                                                                                                                       |
| trace        | Full text trace of agent execution                                                                                                                                                                                                                                                                                                                                                                                    |
| elapsed_time | Execution time in minutes                                                                                                                                                                                                                                                                                                                                                                                             |
| prompt       | The original prompt                                                                                                                                                                                                                                                                                                                                                                                                   |
| trace_json   | Structured JSON trace (when available)                                                                                                                                                                                                                                                                                                                                                                                |
| fold_stats   | Diagnostic metrics from memory folding (list). Includes fold_count (integer or error message string), fold_messages_removed, fold_total_messages_removed, fold_chars_input, fold_summary_chars (effective per-fold summary growth chars), fold_summary_total_chars, fold_summary_delta_chars, fold_trigger_reason, fold_safe_boundary_idx, fold_compression_ratio, fold_parse_success, and fold_summarizer_latency_m. |

|                       |                                                                                |
|-----------------------|--------------------------------------------------------------------------------|
| thread_id             | Resolved LangGraph thread identifier used for this run.                        |
| stop_reason           | Terminal stop reason from agent state (when available).                        |
| budget_state          | Tool-call budget state snapshot from agent state.                              |
| field_status          | Per-field extraction status map from agent state.                              |
| diagnostics           | Runtime diagnostics counters and intervention notes.                           |
| json_repair           | JSON repair/fallback events emitted during execution.                          |
| final_payload         | Canonical terminal payload from agent state (when available).                  |
| terminal_valid        | Logical flag indicating whether a schema-valid terminal payload was emitted.   |
| terminal_payload_hash | Canonical hash of terminal payload from agent state.                           |
| payload_integrity     | Diagnostics about released payload provenance, repair, and truncation markers. |
| completion_gate       | Deterministic outcome-verification report from agent state.                    |
| tokens_used           | Total token count for this invocation (integer, or NA).                        |
| input_tokens          | Input (prompt) token count (integer, or NA).                                   |
| output_tokens         | Output (completion) token count (integer, or NA).                              |
| token_trace           | Per-step token usage trace (list).                                             |

Value

An object of class `asa_response`

---

|            |                                           |
|------------|-------------------------------------------|
| asa_result | <i>Constructor for asa_result Objects</i> |
|------------|-------------------------------------------|

---

Description

Creates an S3 object representing the result of a research task.

Usage

```
asa_result(  
  prompt,  
  message,  
  parsed,  
  raw_output,  
  elapsed_time,  
  status,  
  search_tier = "unknown",  
  parsing_status = NULL,  
  trace_json = "",  
  execution = NULL  
)
```

Arguments

|                |                                                                                                                                                                                                                                                                                                                                                             |
|----------------|-------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| prompt         | The original prompt                                                                                                                                                                                                                                                                                                                                         |
| message        | The agent’s response text                                                                                                                                                                                                                                                                                                                                   |
| parsed         | Parsed output (list or NULL)                                                                                                                                                                                                                                                                                                                                |
| raw_output     | Full agent trace                                                                                                                                                                                                                                                                                                                                            |
| elapsed_time   | Execution time in minutes                                                                                                                                                                                                                                                                                                                                   |
| status         | Status ("success" or "error")                                                                                                                                                                                                                                                                                                                               |
| search_tier    | Which search tier was used ("primp", "selenium", "ddgs", "requests", or "unknown"). Useful for assessing result quality.                                                                                                                                                                                                                                    |
| parsing_status | List with JSON parsing validation info: valid (logical), reason ("ok", "parsing_failed", "not_object", "missing_fields", "null_values", "no_validation"), and missing (character vector of missing/invalid fields).                                                                                                                                         |
| trace_json     | Structured JSON trace (when available)                                                                                                                                                                                                                                                                                                                      |
| execution      | Optional operational metadata list. Common fields include thread_id, stop_reason, status_code, tool budget counters, fold_count, and action/plan metadata extracted from agent state. When created via <a href="#">run_task</a> , selected execution fields are also attached as top-level aliases such as token_stats, fold_stats, plan, and action_ascii. |

Value

An object of class `asa_result`

---

|                   |                                           |
|-------------------|-------------------------------------------|
| benchmark_from_df | Create an ASA Benchmark from a Data Frame |
|-------------------|-------------------------------------------|

---

Description

Converts a tabular benchmark definition into an `asa_benchmark`.

Usage

```
benchmark_from_df(  
  df,  
  prompt_col,  
  expected_col,  
  id_col = NULL,  
  category_col = NULL,  
  output_format_col = NULL,  
  match_col = NULL,  
  run_args_col = NULL,  
  field_matchers_col = NULL,  
  name = "CustomBenchmark",  
  description = NULL,  
  version = "1.0",  
  citation = NULL,  
  metadata = list(),  
  output_format = "text",  
  match = "exact"  
)
```

Arguments

|                    |                                                         |
|--------------------|---------------------------------------------------------|
| df                 | data.frame containing benchmark rows                    |
| prompt_col         | Column containing prompts                               |
| expected_col       | Column containing expected values                       |
| id_col             | Optional column containing task ids                     |
| category_col       | Optional column containing task categories              |
| output_format_col  | Optional column containing per-task output formats      |
| match_col          | Optional column containing per-task match modes         |
| run_args_col       | Optional column containing per-task run arguments       |
| field_matchers_col | Optional column containing per-task field matcher specs |
| name               | Benchmark name (default: "CustomBenchmark")             |
| description        | Optional benchmark description                          |
| version            | Benchmark version string (default: "1.0")               |
| citation           | Optional citation string                                |
| metadata           | Optional benchmark metadata                             |
| output_format      | Default output format when output_format_col is NULL    |
| match              | Default matcher when match_col is NULL                  |

Value

An object of class `asa_benchmark`

---

|                     |                                        |
|---------------------|----------------------------------------|
| benchmark_from_yaml | <i>Load an ASA Benchmark from YAML</i> |
|---------------------|----------------------------------------|

---

Description

Reads a benchmark definition from a YAML file.

Usage

`benchmark_from_yaml(path)`

Arguments

|      |                             |
|------|-----------------------------|
| path | Path to YAML benchmark file |
|------|-----------------------------|

Value

An object of class `asa_benchmark`

---

|               |                                             |
|---------------|---------------------------------------------|
| build_backend | <i>Build the Python Backend Environment</i> |
|---------------|---------------------------------------------|

---

## Description

Creates a conda environment with required Python dependencies for ASA's built-in LangGraph runtime, provider gateways, and search tools.

## Usage

```
build_backend(
  conda_env = NULL,
  conda = "auto",
  python_version = "3.12",
  force = FALSE,
  check_browser = TRUE,
  fix_browser = FALSE,
  chromedriver_dir = NULL
)
```

## Arguments

|                  |                                                                                                                                                                                                                                          |
|------------------|------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| conda_env        | Name of the conda environment. Defaults to the package option <code>asa.default_conda_env</code> (or <code>"asa_env"</code> if unset).                                                                                                   |
| conda            | Path to conda executable (default: <code>"auto"</code> )                                                                                                                                                                                 |
| python_version   | Python version to use (default: <code>"3.12"</code> )                                                                                                                                                                                    |
| force            | If TRUE, delete and recreate the conda environment if it already exists.                                                                                                                                                                 |
| check_browser    | If TRUE, performs a best-effort check for a system Chrome/Chromium binary and chromedriver, warning when major versions are incompatible. Set to FALSE to skip this check.                                                               |
| fix_browser      | If TRUE and <code>check_browser = TRUE</code> , attempt to download a matching chromedriver for the detected Chrome version and prepend it to PATH for the current R session. The driver is stored under <code>chromedriver_dir</code> . |
| chromedriver_dir | Optional directory to store downloaded chromedriver binaries. Defaults to <code>~/ .asa/chromedriver</code> when <code>fix_browser = TRUE</code> .                                                                                       |

## Details

This function creates a new conda environment and installs the following Python packages:

- `langchain_groq`, `langchain_community`, `langchain_openai`, `langchain_google_genai`
- `langgraph`
- `ddgs` (DuckDuckGo search)
- `selenium`, `primp` (browser automation)
- `undetected-chromedriver` (stealth Chrome)
- `beautifulsoup4`, `curl_cffi`, `requests`
- `langextract` (schema extraction from webpage text)

- fake\_headers, httpx
- stem (Tor control)
- pysocks, socksio (proxy support)

Python compatibility note: Current LangChain releases still import parts of pydantic.v1, which emits warnings on Python 3.14+. Prefer Python 3.12 (or 3.13) until upstream dependencies remove that compatibility layer.

## Value

Invisibly returns NULL; called for side effects.

## Examples

```
## Not run:
# Create the default environment
build_backend()

# Rebuild from scratch
build_backend(force = TRUE)

# Create with a custom name
build_backend(conda_env = "my_asa_env")

## End(Not run)
```

---

|              |                                          |
|--------------|------------------------------------------|
| build_prompt | <i>Build a Task Prompt from Template</i> |
|--------------|------------------------------------------|

---

## Description

Creates a formatted prompt by substituting variables into a template.

## Usage

```
build_prompt(template, ...)
```

## Arguments

|          |                                                                  |
|----------|------------------------------------------------------------------|
| template | A character string with placeholders in the form {variable_name} |
| ...      | Named arguments to substitute into the template                  |

## Value

A formatted prompt string

## Examples

```
## Not run:
prompt <- build_prompt(
  template = "Find information about {{name}} in {{country}} during {{year}}",
  name = "Marie Curie",
  country = "France",
  year = 1903
)

## End(Not run)
```

---

|               |                                          |
|---------------|------------------------------------------|
| check_backend | <i>Check Python Backend Availability</i> |
|---------------|------------------------------------------|

---

## Description

Checks whether the configured Conda environment exists and can import the Python modules required by the ASA backend, without initializing reticulate's Python session for the current R process.

## Usage

```
check_backend(conda_env = NULL, conda = "auto", strict = FALSE)
```

## Arguments

|           |                                                                                                                                                            |
|-----------|------------------------------------------------------------------------------------------------------------------------------------------------------------|
| conda_env | Name of the Conda environment. Defaults to the package option <code>asa.default_conda_env</code> (or <code>"asa_env"</code> if unset).                     |
| conda     | Path to conda executable (default: <code>"auto"</code> )                                                                                                   |
| strict    | If TRUE, also require optional packages used for Tor control and stealth Chrome support (e.g., <code>stem</code> , <code>undetected_chromedriver</code> ). |

## Value

A list with components:

- `available`: Logical, TRUE if environment is ready
- `conda_env`: Name of the environment checked
- `python_version`: Python version if available
- `missing_packages`: Character vector of missing required packages (if any)
- `missing_optional_packages`: Character vector of missing optional packages (if any)
- `strict`: Logical, whether optional packages were required

**Examples**

```
## Not run:
status <- check_backend()
if (!status$available) {
  build_backend()
}

## End(Not run)
```

configure\_search

*Configure Python Search Parameters***Description**

Sets global configuration values for the Python search module. These values control timeouts, retry behavior, and result limits.

**Usage**

```
configure_search(
  max_results = NULL,
  timeout = NULL,
  max_retries = NULL,
  retry_delay = NULL,
  backoff_multiplier = NULL,
  captcha_backoff_base = NULL,
  page_load_wait = NULL,
  inter_search_delay = NULL,
  conda_env = NULL,
  selenium_browser_preference = NULL
)
```

**Arguments**

|                                          |                                                                                                                                        |
|------------------------------------------|----------------------------------------------------------------------------------------------------------------------------------------|
| <code>max_results</code>                 | Maximum number of search results to return (default: 10)                                                                               |
| <code>timeout</code>                     | HTTP request timeout in seconds (default: 30)                                                                                          |
| <code>max_retries</code>                 | Maximum retry attempts on failure (default: 3)                                                                                         |
| <code>retry_delay</code>                 | Initial delay between retries in seconds (default: 2)                                                                                  |
| <code>backoff_multiplier</code>          | Multiplier for exponential backoff (default: 1.5)                                                                                      |
| <code>captcha_backoff_base</code>        | Base multiplier for CAPTCHA backoff (default: 3)                                                                                       |
| <code>page_load_wait</code>              | Wait time after page load in seconds (default: 2)                                                                                      |
| <code>inter_search_delay</code>          | Delay between consecutive searches in seconds (default: 1.5)                                                                           |
| <code>conda_env</code>                   | Name of the conda environment. Defaults to the package option <code>asa.default_conda_env</code> (or <code>"asa_env"</code> if unset). |
| <code>selenium_browser_preference</code> | Selenium browser engine preference order. One of <code>"firefox_first"</code> (default) or <code>"chrome_first"</code> .               |



**Value**

Invisibly returns a list with the current configuration

**Examples**

```
## Not run:
# Increase timeout for slow connections
configure_search(timeout = 30, max_retries = 5)

# Get more results
configure_search(max_results = 20)

# Add delay between searches to avoid rate limiting
configure_search(inter_search_delay = 2.0)

## End(Not run)
```

---

```
configure_search_logging
```

*Configure Python Search Logging Level*

---

**Description**

Sets the logging level for the Python search module. This controls how much diagnostic output is produced during web searches.

**Usage**

```
configure_search_logging(level = "WARNING", conda_env = NULL)
```

**Arguments**

|           |                                                                                                                           |
|-----------|---------------------------------------------------------------------------------------------------------------------------|
| level     | Log level: "DEBUG", "INFO", "WARNING" (default), "ERROR", or "CRITICAL"                                                   |
| conda_env | Name of the conda environment. Defaults to the package option <code>asa.default_conda_env</code> (or "asa_env" if unset). |

**Details**

Log levels from most to least verbose:

- DEBUG: Detailed diagnostic information for debugging
- INFO: General operational information
- WARNING: Indicates something unexpected but not an error (default)
- ERROR: Serious problems that prevented an operation
- CRITICAL: Very serious errors

**Value**

Invisibly returns the current logging level

**Examples**

```
## Not run:
# Enable verbose debugging output
configure_search_logging("DEBUG")

# Run a search (will show detailed logs)
result <- run_task("What is the population of Tokyo?", agent = agent)

# Disable verbose output
configure_search_logging("WARNING")

## End(Not run)
```

---

|                    |                                                |
|--------------------|------------------------------------------------|
| configure_temporal | <i>Configure Temporal Filtering for Search</i> |
|--------------------|------------------------------------------------|

---

**Description**

Sets or clears temporal filtering on the DuckDuckGo search tool. This affects all subsequent searches until changed or cleared.

**Usage**

```
configure_temporal(time_filter = NULL)
```

**Arguments**

|             |                                                                                                    |
|-------------|----------------------------------------------------------------------------------------------------|
| time_filter | DuckDuckGo time filter: "d" (day), "w" (week), "m" (month), "y" (year), or NULL/NA/"none" to clear |
|-------------|----------------------------------------------------------------------------------------------------|

**Details**

This function modifies the search tool's time parameter, which is passed to DuckDuckGo as the df parameter. The filter restricts results to content indexed within the specified time period.

Note: This only affects DuckDuckGo searches. For Wikidata queries with temporal filtering, use `asa_enumerate()` with its temporal parameter.

**Value**

Invisibly returns the previous time filter setting

**Time Filter Values**

- "d": Past 24 hours (day)
- "w": Past 7 days (week)
- "m": Past 30 days (month)
- "y": Past 365 days (year)
- NULL, NA, or "none": No time restriction (default)

**See Also**

[run\\_task](#), [asa\\_enumerate](#)

**Examples**

```
## Not run:
# Restrict to past year
configure_temporal("y")
result <- run_task("Find recent AI breakthroughs", agent = agent)

# Clear temporal filter
configure_temporal(NULL)

# Past week only
configure_temporal("w")

## End(Not run)
```

---

```
configure_tor_registry
```

*Configure Tor Exit Registry*

---

**Description**

Sets up the shared Tor exit health registry used by the Python search stack to avoid reusing tainted or overused exit nodes.

**Usage**

```
configure_tor_registry(
  registry_path = NULL,
  enable = ASA_TOR_REGISTRY_ENABLED,
  bad_ttl = ASA_TOR_BAD_TTL,
  good_ttl = ASA_TOR_GOOD_TTL,
  overuse_threshold = ASA_TOR_OVERUSE_THRESHOLD,
  overuse_decay = ASA_TOR_OVERUSE_DECAY,
  max_rotation_attempts = ASA_TOR_MAX_ROTATION_ATTEMPTS,
  ip_cache_ttl = ASA_TOR_IP_CACHE_TTL,
  conda_env = NULL
)
```

**Arguments**

|                   |                                                                  |
|-------------------|------------------------------------------------------------------|
| registry_path     | Path to the SQLite registry file (default: user cache).          |
| enable            | Enable the registry (set FALSE to disable tracking).             |
| bad_ttl           | Seconds to keep a bad/tainted exit before reuse.                 |
| good_ttl          | Seconds to treat an exit as good before refreshing.              |
| overuse_threshold | Maximum recent uses before a good exit is treated as overloaded. |

|                       |                                                                                                                                                       |
|-----------------------|-------------------------------------------------------------------------------------------------------------------------------------------------------|
| overuse_decay         | Window (seconds) for counting recent uses before decay.                                                                                               |
| max_rotation_attempts | Maximum rotations to find a clean exit.                                                                                                               |
| ip_cache_ttl          | Seconds to cache exit IP lookups.                                                                                                                     |
| conda_env             | Conda environment name for the Python module. Defaults to the package option <code>asa.default_conda_env</code> (or <code>"asa_env"</code> if unset). |

**Value**

Invisibly returns a list of the configured values (or NULL on error).

---

|                  |                                       |
|------------------|---------------------------------------|
| create_benchmark | <i>Create an ASA Benchmark Object</i> |
|------------------|---------------------------------------|

---

**Description**

Creates a structured benchmark specification for agent evaluation.

**Usage**

```
create_benchmark(  
    name,  
    tasks,  
    description = NULL,  
    version = "1.0",  
    citation = NULL,  
    metadata = list()  
)
```

**Arguments**

|             |                                                      |
|-------------|------------------------------------------------------|
| name        | Benchmark name                                       |
| tasks       | List of benchmark task definitions                   |
| description | Optional benchmark description                       |
| version     | Benchmark version string (default: "1.0")            |
| citation    | Optional citation string                             |
| metadata    | Optional named list of additional benchmark metadata |

**Value**

An object of class `asa_benchmark`

---

|                |                                                  |
|----------------|--------------------------------------------------|
| evaluate_agent | <i>Evaluate an ASA Agent Against a Benchmark</i> |
|----------------|--------------------------------------------------|

---

**Description**

Runs a benchmark suite against an ASA agent using deterministic local scoring.

**Usage**

```
evaluate_agent(
  benchmark,
  agent = NULL,
  config = NULL,
  runs = 1L,
  fail_fast = FALSE,
  keep_outputs = FALSE,
  verbose = TRUE
)
```

**Arguments**

|              |                                                                            |
|--------------|----------------------------------------------------------------------------|
| benchmark    | An asa_benchmark object or built-in benchmark name                         |
| agent        | Optional initialized asa_agent                                             |
| config       | Optional asa_config used to initialize an agent when agent is not supplied |
| runs         | Number of repeated benchmark runs (default: 1)                             |
| fail_fast    | If TRUE, stop evaluation on first execution failure                        |
| keep_outputs | If TRUE, retain raw asa_result objects in the returned evaluation result   |
| verbose      | Print progress messages                                                    |

**Value**

An object of class asa\_evaluation\_result

---

|                       |                                                  |
|-----------------------|--------------------------------------------------|
| extract_agent_results | <i>Extract Structured Data from Agent Traces</i> |
|-----------------------|--------------------------------------------------|

---

**Description**

Parses raw agent output to extract search snippets, Wikipedia content, URLs, JSON data, and search tier information. This is the main function for post-processing agent traces.

**Usage**

```
extract_agent_results(raw_output)
```

**Arguments**

|            |                                                                                                                     |
|------------|---------------------------------------------------------------------------------------------------------------------|
| raw_output | Raw trace string from asa_result\$raw_output, or a structured JSON trace (asa_trace_v1) from asa_result\$trace_json |
|------------|---------------------------------------------------------------------------------------------------------------------|

**Value**

A list with components:

- search\_snippets: Character vector of search result content
- search\_urls: Character vector of URLs from search results
- wikipedia\_snippets: Character vector of Wikipedia content
- json\_data: Extracted JSON data (canonical by default; see option `asa.enable_trace_field_mining`)
- json\_data\_canonical: Canonical terminal JSON extracted from trace
- json\_data\_inferred: Optional inferred JSON with trace-mining enabled
- search\_tiers: Character vector of unique search tiers used (e.g., "primp", "selenium", "ddgs", "requests")

**Examples**

```
## Not run:
result <- run_task("Who is the president of France?", agent = agent)
trace <- if (!is.null(result$trace_json) && nzchar(result$trace_json)) {
  result$trace_json
} else {
  result$raw_output
}
extracted <- extract_agent_results(trace)
print(extracted$search_snippets)
print(extracted$search_tiers) # Shows which search tier was used

## End(Not run)
```

---

```
extract_search_snippets
```

*Extract Search Snippets by Source Number*

---

**Description**

Extracts content from Search tool messages in the agent trace.

**Usage**

```
extract_search_snippets(text, source = NULL)
```

**Arguments**

|        |                                                                                                                                                                                      |
|--------|--------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| text   | Structured JSON trace ( <code>asa_trace_v1</code> )                                                                                                                                  |
| source | Optional integer vector of source numbers to return (1-indexed). If NULL (default), returns a character vector for all sources encountered (with empty strings for missing indices). |

**Value**

Character vector of search snippets, ordered by source number

## Examples

```
## Not run:
snippets <- extract_search_snippets(response$trace)

## End(Not run)
```

---

|                      |                                        |
|----------------------|----------------------------------------|
| extract_search_tiers | <i>Extract Search Tier Information</i> |
|----------------------|----------------------------------------|

---

## Description

Extracts which search tier was used from the agent trace. The search module uses a multi-tier fallback system:

- primp: Fast HTTP client with browser impersonation (Tier 0)
- selenium: Headless browser for JS-rendered content (Tier 1)
- ddgs: Standard DDGS Python library (Tier 2)
- requests: Raw POST to DuckDuckGo HTML endpoint (Tier 3)

## Usage

```
extract_search_tiers(text)
```

## Arguments

|      |                                      |
|------|--------------------------------------|
| text | Structured JSON trace (asa_trace_v1) |
|------|--------------------------------------|

## Value

Character vector of unique tier names encountered (e.g., "primp", "selenium", "ddgs", "requests")

## Examples

```
## Not run:
tiers <- extract_search_tiers(response$trace)
print(tiers) # e.g., "primp"

## End(Not run)
```

---

|              |                                      |
|--------------|--------------------------------------|
| extract_urls | <i>Extract URLs by Source Number</i> |
|--------------|--------------------------------------|

---

**Description**

Extracts URLs from Search tool messages in the agent trace.

**Usage**

```
extract_urls(text, source = NULL)
```

**Arguments**

|        |                                                                                                                                                                                      |
|--------|--------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| text   | Structured JSON trace (asa_trace_v1)                                                                                                                                                 |
| source | Optional integer vector of source numbers to return (1-indexed). If NULL (default), returns a character vector for all sources encountered (with empty strings for missing indices). |

**Value**

Character vector of URLs, ordered by source number

**Examples**

```
## Not run:  
urls <- extract_urls(response$trace)  
  
## End(Not run)
```

---

|                           |                                  |
|---------------------------|----------------------------------|
| extract_wikipedia_content | <i>Extract Wikipedia Content</i> |
|---------------------------|----------------------------------|

---

**Description**

Extracts content from Wikipedia tool messages in the agent trace.

**Usage**

```
extract_wikipedia_content(text)
```

**Arguments**

|      |                                      |
|------|--------------------------------------|
| text | Structured JSON trace (asa_trace_v1) |
|------|--------------------------------------|

**Value**

Character vector of Wikipedia snippets



**Examples**

```
## Not run:
wiki <- extract_wikipedia_content(response$trace)

## End(Not run)
```

get\_agent

*Get the Current Agent***Description**

Returns the currently initialized agent, or NULL if not initialized.

**Usage**

```
get_agent()
```

**Value**

An asa\_agent object or NULL

**Examples**

```
## Not run:
agent <- get_agent()
if (is.null(agent)) {
  agent <- initialize_agent()
}

## End(Not run)
```

get\_tor\_ip

*Get External IP via Tor***Description**

Retrieves the external IP address as seen through Tor proxy.

**Usage**

```
get_tor_ip(proxy = "socks5h://127.0.0.1:9050", timeout = 30L)
```

**Arguments**

|         |                                                                                                                       |
|---------|-----------------------------------------------------------------------------------------------------------------------|
| proxy   | Tor proxy URL (e.g., "socks5h://127.0.0.1:9050" for default, or "socks5h://127.0.0.1:9055" for instance on port 9055) |
| timeout | Timeout in seconds (default: 30). Useful for parallel workloads where some Tor exits may be slow.                     |

**Value**

IP address string or NA on failure

**Examples**

```
## Not run:
# Default Tor instance
ip <- get_tor_ip()
message("Current Tor IP: ", ip)

# Check specific Tor instance (e.g., for parallel jobs)
ip <- get_tor_ip(proxy = "socks5h://127.0.0.1:9055")

## End(Not run)
```

---

initialize\_agent

*Initialize the ASA Search Agent*


---

**Description**

Initializes the Python environment and creates an ASA research agent runtime with search tools (Wikipedia, DuckDuckGo). The default runtime is OpenCode; the built-in LangGraph pipeline and free-code runtime remain available. The agent can use multiple LLM backends and supports DeepAgent-style memory folding.

**Usage**

```
initialize_agent(
  agent_backend = NULL,
  backend = NULL,
  model = NULL,
  conda_env = NULL,
  proxy = NA,
  use_browser = ASA_DEFAULT_USE_BROWSER,
  search = NULL,
  use_memory_folding = ASA_DEFAULT_MEMORY_FOLDING,
  memory_threshold = ASA_DEFAULT_MEMORY_THRESHOLD,
  memory_keep_recent = ASA_DEFAULT_MEMORY_KEEP_RECENT,
  fold_char_budget = ASA_DEFAULT_FOLD_CHAR_BUDGET,
  use_observational_memory = ASA_DEFAULT_USE_OBSERVATIONAL_MEMORY,
  om_observation_token_budget = ASA_DEFAULT_OM_OBSERVATION_TOKENS,
  om_reflection_token_budget = ASA_DEFAULT_OM_REFLECTION_TOKENS,
  om_buffer_tokens = ASA_DEFAULT_OM_BUFFER_TOKENS,
  om_buffer_activation = ASA_DEFAULT_OM_BUFFER_ACTIVATION,
  om_block_after = ASA_DEFAULT_OM_BLOCK_AFTER,
  om_async_prebuffer = ASA_DEFAULT_OM_ASYNC_PREBUFFER,
  om_cross_thread_memory = ASA_DEFAULT_OM_CROSS_THREAD_MEMORY,
  rate_limit = ASA_DEFAULT_RATE_LIMIT,
  timeout = ASA_DEFAULT_TIMEOUT,
  tor = tor_options(),
```

```

    verbose = TRUE,
    recursion_limit = NULL
)

```

## Arguments

|                             |                                                                                                                                                                                                                                                                                                                            |
|-----------------------------|----------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| agent_backend               | Agent runtime backend. Use "opencode" (default) to run the OpenCode CLI behind ASA-managed provider and search wrappers, "agent" for the built-in ASA LangGraph pipeline, or "free-code" to run the external free-code agent behind the same ASA-managed provider and search wrappers.                                     |
| backend                     | LLM backend to use. One of: "openai", "azure-openai", "groq", "xai", "gemini", "exo", "ollama", "openrouter", "anthropic", "bedrock". Ollama also supports shorthand like "ollama-lfm2:24b-a2b" when model is omitted.                                                                                                     |
| model                       | Model identifier (e.g., "gpt-4.1-mini", "llama-3.3-70b-versatile")                                                                                                                                                                                                                                                         |
| conda_env                   | Name of the conda environment with Python dependencies                                                                                                                                                                                                                                                                     |
| proxy                       | Proxy URL for search tools. Use NA (default) to auto-detect from environment variables (ASA_PROXY, HTTP_PROXY, HTTPS_PROXY). Use NULL to disable proxying. Use ASA_DEFAULT_PROXY (or "socks5h://127.0.0.1:9050") to route searches through Tor.                                                                            |
| use_browser                 | Enable Selenium browser tier for DuckDuckGo search (default: TRUE). Set to FALSE to avoid Chromedriver requirements.                                                                                                                                                                                                       |
| search                      | Optional search options (created with <a href="#">search_options</a> ) used as defaults for <a href="#">run_task/asa_enumerate</a> when no asa_config is provided. This is useful for setting OpenWebpage tuning knobs (e.g., webpage_max_chars, webpage_max_chunks, webpage_chunk_chars) once when initializing an agent. |
| use_memory_folding          | Enable DeepAgent-style memory compression (default: TRUE)                                                                                                                                                                                                                                                                  |
| memory_threshold            | Number of messages before folding triggers (default: 10; message-count back-stop alongside fold_char_budget)                                                                                                                                                                                                               |
| memory_keep_recent          | Number of recent exchanges to preserve after folding (default: 4). An exchange is a user turn plus the assistant response, including any tool calls and tool outputs.                                                                                                                                                      |
| fold_char_budget            | Total character budget across all messages before memory folding triggers (default: 30000). Lower values fold more aggressively; higher values preserve more raw conversation.                                                                                                                                             |
| use_observational_memory    | Enable observational memory pipeline.                                                                                                                                                                                                                                                                                      |
| om_observation_token_budget | Token budget for stored observation messages.                                                                                                                                                                                                                                                                              |
| om_reflection_token_budget  | Token budget for reflection summaries.                                                                                                                                                                                                                                                                                     |
| om_buffer_tokens            | Token budget for buffered context before reflection.                                                                                                                                                                                                                                                                       |
| om_buffer_activation        | Fraction of buffer usage that triggers reflection.                                                                                                                                                                                                                                                                         |
| om_block_after              | Fraction of context usage where buffering blocks generation.                                                                                                                                                                                                                                                               |

|                        |                                                                                                                                                                                                                                                                                                |
|------------------------|------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| om_async_prebuffer     | Whether to pre-buffer observations asynchronously.                                                                                                                                                                                                                                             |
| om_cross_thread_memory | Whether observational memory is shared across threads.                                                                                                                                                                                                                                         |
| rate_limit             | Requests per second for rate limiting (default: 0.1)                                                                                                                                                                                                                                           |
| timeout                | Request timeout in seconds (default: 120)                                                                                                                                                                                                                                                      |
| tor                    | Tor registry options from <a href="#">tor_options</a> . Disable shared tracking by setting <code>dirty_tor_exists = FALSE</code> .                                                                                                                                                             |
| verbose                | Print status messages (default: TRUE)                                                                                                                                                                                                                                                          |
| recursion_limit        | Optional default maximum number of agent steps to store in the initialized agent. When NULL (default), step limits fall back to mode-specific defaults at run time (memory folding: 100; standard agent: 20). Per-call <code>run_task(..., recursion_limit = ...)</code> overrides this value. |

## Details

The agent is created with these tools:

- Wikipedia: For looking up encyclopedic information
- DuckDuckGo Search: For web searches with a 4-tier fallback system (PRIMP -> Selenium -> DDGS library -> raw requests)
- OpenWebpage: (Optional) Fetch and extract readable text from a public webpage URL. This capability is disabled by default and must be enabled per-call (e.g., via `run_task(..., allow_read_webpages = TRUE)`).

Memory folding (enabled by default) compresses older messages into a summary to manage context length in long conversations, following the DeepAgent paper.

## Value

An object of class `asa_agent` containing the initialized agent and configuration.

## API Keys

The following environment variables should be set based on your backend:

- OpenAI: `OPENAI_API_KEY`
- Azure OpenAI: `AZURE_OPENAI_API_KEY` plus `AZURE_OPENAI_ENDPOINT` or `AZURE_OPENAI_API_BASE`; `model` is the Azure deployment name.
- Groq: `GROQ_API_KEY`
- xAI: `XAI_API_KEY`
- Gemini: `GOOGLE_API_KEY` (or `GEMINI_API_KEY`)
- Anthropic: `ANTHROPIC_API_KEY`
- Bedrock: `AWS_ACCESS_KEY_ID` + `AWS_SECRET_ACCESS_KEY` (optionally `AWS_DEFAULT_REGION`, defaults to "us-east-1")
- OpenRouter: `OPENROUTER_API_KEY`
- Exo: no API key required (local server)
- Ollama: no API key required (local server)

## OpenRouter Models

When using the "openrouter" backend, model names must be in provider/model-name format.  
Examples:

- "openai/gpt-4o"
- "anthropic/claude-3-sonnet"
- "google/gemma-2-9b-it:free"
- "meta-llama/llama-3-70b-instruct"

See <https://openrouter.ai/models> for available models.

## See Also

[run\\_task](#), [run\\_task\\_batch](#)

## Examples

```
## Not run:
# Initialize with OpenAI
agent <- initialize_agent(
  backend = "openai",
  model = "gpt-4.1-mini"
)

# Initialize with Groq and custom settings
agent <- initialize_agent(
  backend = "groq",
  model = "llama-3.3-70b-versatile",
  use_memory_folding = FALSE,
  proxy = NULL # No Tor proxy
)

# Initialize with OpenRouter (access to 100+ models)
agent <- initialize_agent(
  backend = "openrouter",
  model = "anthropic/claude-3-sonnet" # Note: provider/model format
)

# Initialize with a local Ollama model
agent <- initialize_agent(
  backend = "ollama",
  model = "l1m2:24b-a2b"
)

## End(Not run)
```

---

is\_tor\_running

---

*Check if Tor is Running*


---

## Description

Checks if Tor is running and accessible on the default port.

Usage

```
is_tor_running(port = 9050L)
```

Arguments

port                      Port number (default: 9050)

Value

Logical indicating if Tor appears to be running

Examples

```
## Not run:
if (!is_tor_running()) {
  message("Start Tor with: brew services start tor")
}

## End(Not run)
```

---

|                 |                                     |
|-----------------|-------------------------------------|
| list_benchmarks | <i>List Built-In ASA Benchmarks</i> |
|-----------------|-------------------------------------|

---

Description

Lists benchmark definitions bundled with the package.

Usage

```
list_benchmarks()
```

Value

data.frame of benchmark metadata

---

|                |                                      |
|----------------|--------------------------------------|
| load_benchmark | <i>Load a Built-In ASA Benchmark</i> |
|----------------|--------------------------------------|

---

Description

Loads a benchmark bundled under inst/benchmarks/.

Usage

```
load_benchmark(name)
```

Arguments

name                      Benchmark name

**Value**

An object of class asa\_benchmark

---

|                 |                                           |
|-----------------|-------------------------------------------|
| print.asa_agent | <i>Print Method for asa_agent Objects</i> |
|-----------------|-------------------------------------------|

---

**Description**

Print Method for asa\_agent Objects

**Usage**

```
## S3 method for class 'asa_agent'  
print(x, ...)
```

**Arguments**

|     |                                |
|-----|--------------------------------|
| x   | An asa_agent object            |
| ... | Additional arguments (ignored) |

**Value**

Invisibly returns the object

---

|                        |                                                  |
|------------------------|--------------------------------------------------|
| print.asa_audit_result | <i>Print Method for asa_audit_result Objects</i> |
|------------------------|--------------------------------------------------|

---

**Description**

Print Method for asa\_audit\_result Objects

**Usage**

```
## S3 method for class 'asa_audit_result'  
print(x, n = 6, ...)
```

**Arguments**

|     |                                             |
|-----|---------------------------------------------|
| x   | An asa_audit_result object                  |
| n   | Number of data rows to preview (default: 6) |
| ... | Additional arguments (ignored)              |

**Value**

Invisibly returns the object

---

|                     |                                               |
|---------------------|-----------------------------------------------|
| print.asa_benchmark | <i>Print Method for asa_benchmark Objects</i> |
|---------------------|-----------------------------------------------|

---

**Description**

Print Method for asa\_benchmark Objects

**Usage**

```
## S3 method for class 'asa_benchmark'  
print(x, ...)
```

**Arguments**

|     |                                |
|-----|--------------------------------|
| x   | An asa_benchmark object        |
| ... | Additional arguments (ignored) |

**Value**

Invisibly returns the benchmark object

---

|                  |                                            |
|------------------|--------------------------------------------|
| print.asa_config | <i>Print Method for asa_config Objects</i> |
|------------------|--------------------------------------------|

---

**Description**

Print Method for asa\_config Objects

**Usage**

```
## S3 method for class 'asa_config'  
print(x, ...)
```

**Arguments**

|     |                                |
|-----|--------------------------------|
| x   | An asa_config object           |
| ... | Additional arguments (ignored) |

**Value**

Invisibly returns the object



---

```
print.asa_enumerate_result
```

*Print Method for asa\_enumerate\_result Objects*

---

### Description

Print Method for asa\_enumerate\_result Objects

### Usage

```
## S3 method for class 'asa_enumerate_result'  
print(x, n = 6, ...)
```

### Arguments

|     |                                             |
|-----|---------------------------------------------|
| x   | An asa_enumerate_result object              |
| n   | Number of data rows to preview (default: 6) |
| ... | Additional arguments (ignored)              |

### Value

Invisibly returns the object

---

```
print.asa_evaluation_result
```

*Print Method for asa\_evaluation\_result Objects*

---

### Description

Print Method for asa\_evaluation\_result Objects

### Usage

```
## S3 method for class 'asa_evaluation_result'  
print(x, ...)
```

### Arguments

|     |                                 |
|-----|---------------------------------|
| x   | An asa_evaluation_result object |
| ... | Additional arguments (ignored)  |

### Value

Invisibly returns the evaluation result

---

|                    |                                              |
|--------------------|----------------------------------------------|
| print.asa_response | <i>Print Method for asa_response Objects</i> |
|--------------------|----------------------------------------------|

---

**Description**

Print Method for asa\_response Objects

**Usage**

```
## S3 method for class 'asa_response'  
print(x, ...)
```

**Arguments**

|     |                                |
|-----|--------------------------------|
| x   | An asa_response object         |
| ... | Additional arguments (ignored) |

**Value**

Invisibly returns the object

---

|                  |                                            |
|------------------|--------------------------------------------|
| print.asa_result | <i>Print Method for asa_result Objects</i> |
|------------------|--------------------------------------------|

---

**Description**

Print Method for asa\_result Objects

**Usage**

```
## S3 method for class 'asa_result'  
print(x, ...)
```

**Arguments**

|     |                                |
|-----|--------------------------------|
| x   | An asa_result object           |
| ... | Additional arguments (ignored) |

**Value**

Invisibly returns the object

---

|                  |                                            |
|------------------|--------------------------------------------|
| print.asa_search | <i>Print Method for asa_search Objects</i> |
|------------------|--------------------------------------------|

---

**Description**

Print Method for asa\_search Objects

**Usage**

```
## S3 method for class 'asa_search'  
print(x, ...)
```

**Arguments**

|     |                                |
|-----|--------------------------------|
| x   | An asa_search object           |
| ... | Additional arguments (ignored) |

---

|                    |                                              |
|--------------------|----------------------------------------------|
| print.asa_temporal | <i>Print Method for asa_temporal Objects</i> |
|--------------------|----------------------------------------------|

---

**Description**

Print Method for asa\_temporal Objects

**Usage**

```
## S3 method for class 'asa_temporal'  
print(x, ...)
```

**Arguments**

|     |                                |
|-----|--------------------------------|
| x   | An asa_temporal object         |
| ... | Additional arguments (ignored) |

**Value**

Invisibly returns the object

---

|               |                                         |
|---------------|-----------------------------------------|
| print.asa_tor | <i>Print Method for asa_tor Objects</i> |
|---------------|-----------------------------------------|

---

**Description**

Print Method for asa\_tor Objects

**Usage**

```
## S3 method for class 'asa_tor'  
print(x, ...)
```

**Arguments**

|     |                                |
|-----|--------------------------------|
| x   | An asa_tor object              |
| ... | Additional arguments (ignored) |

**Value**

Invisibly returns the object

---

|                 |                                       |
|-----------------|---------------------------------------|
| process_outputs | <i>Process Multiple Agent Outputs</i> |
|-----------------|---------------------------------------|

---

**Description**

Processes a data frame of raw agent outputs, extracting structured data.

**Usage**

```
process_outputs(df, parallel = FALSE, workers = 10L)
```

**Arguments**

|          |                                                               |
|----------|---------------------------------------------------------------|
| df       | Data frame with a 'raw_output' column containing agent traces |
| parallel | Use parallel processing                                       |
| workers  | Number of workers                                             |

**Value**

The input data frame with additional extracted columns: search\_count, wiki\_count, and any JSON fields found

---

|             |                        |
|-------------|------------------------|
| reset_agent | <i>Reset the Agent</i> |
|-------------|------------------------|

---

**Description**

Clears the initialized agent state, forcing reinitialization on next use. Also closes any open HTTP clients to prevent resource leaks.

**Usage**

```
reset_agent()
```

**Value**

Invisibly returns NULL

---

|                    |                                                    |
|--------------------|----------------------------------------------------|
| rotate_tor_circuit | <i>Rotate Tor Circuit (R-side, daemon restart)</i> |
|--------------------|----------------------------------------------------|

---

**Description**

Requests a new Tor circuit by restarting the Tor service or sending SIGHUP.

**Usage**

```
rotate_tor_circuit(
    method = c("signal", "brew", "systemctl"),
    wait = 12L,
    pid = NULL
)
```

**Arguments**

|        |                                                                                                                            |
|--------|----------------------------------------------------------------------------------------------------------------------------|
| method | Method to restart: "brew" (macOS), "systemctl" (Linux), or "signal"                                                        |
| wait   | Seconds to wait for new circuit (default: 12)                                                                              |
| pid    | Optional PID of specific Tor process (only used with method="signal"). If NULL (default), finds the Tor process via pgrep. |

**Details**

**MEDIUM FIX:** This function restarts the entire Tor daemon, which kills ALL circuits and affects parallel execution. For production use, prefer the Python-side control port rotation which sends SIGNAL NEWNYM to get a new circuit without restarting the daemon.

For parallel Tor setups with multiple instances, consider using Tor's built-in circuit rotation via MaxCircuitDirtiness and NewCircuitPeriod config options instead of this function.

**Value**

Invisibly returns TRUE on success, FALSE on failure

**Note**

The "brew" and "systemctl" methods restart the entire Tor daemon and should only be used as a last resort for recovery. The "signal" method is preferred but still affects all circuits on the process.

**Examples**

```
## Not run:
# Preferred: Use Python-side control port rotation (via run_task/asa_enumerate)
# This R function is for manual recovery only

# Send SIGHUP to Tor process (least disruptive)
rotate_tor_circuit(method = "signal")

# macOS with Homebrew (restarts daemon - use sparingly)
rotate_tor_circuit(method = "brew")

# Linux with systemd (restarts daemon - use sparingly)
rotate_tor_circuit(method = "systemctl")

## End(Not run)
```

---

|                 |                                                       |
|-----------------|-------------------------------------------------------|
| run_direct_task | <i>Run a Direct Provider Task Without ASA Tooling</i> |
|-----------------|-------------------------------------------------------|

---

**Description**

Invokes the configured provider directly using the underlying chat model attached to an initialized agent. This bypasses the ASA agent loop, tools, search, and temporal prompt augmentation while preserving the same backend and model selection.

**Usage**

```
run_direct_task(
  prompt,
  output_format = "text",
  config = NULL,
  agent = NULL,
  verbose = FALSE
)
```

**Arguments**

|               |                                                                                                                                                                                                                                                                                                      |
|---------------|------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| prompt        | The prompt or question to send directly to the configured LLM                                                                                                                                                                                                                                        |
| output_format | Expected output format. One of: <ul style="list-style-type: none"><li>"text": Returns response text (default)</li><li>"json": Parse response as JSON</li><li>"raw": Include the raw provider response object for debugging</li><li>Character vector: Extract specific fields from response</li></ul> |
| config        | An asa_config object for unified configuration, or NULL to use the currently initialized agent.                                                                                                                                                                                                      |

|         |                                                                                                          |
|---------|----------------------------------------------------------------------------------------------------------|
| agent   | An asa_agent object from <a href="#">initialize_agent</a> , or NULL to initialize/reuse one from config. |
| verbose | Print progress messages (default: FALSE)                                                                 |

### Value

An asa\_result object with execution\$mode = "provider\_direct" and search\_tier = "none".

---

|          |                                             |
|----------|---------------------------------------------|
| run_task | <i>Run a Structured Task with the Agent</i> |
|----------|---------------------------------------------|

---

### Description

Executes a research task using the AI search agent with a structured prompt and returns parsed results. This is the primary function for running agent tasks.

### Usage

```
run_task(
  prompt,
  output_format = "text",
  temporal = NULL,
  config = NULL,
  agent = NULL,
  expected_fields = NULL,
  expected_schema = NULL,
  thread_id = NULL,
  field_status = NULL,
  budget_state = NULL,
  search_budget_limit = NULL,
  unknown_after_searches = NULL,
  finalize_on_all_fields_resolved = NULL,
  field_rules = NULL,
  source_policy = NULL,
  retry_policy = NULL,
  finalization_policy = NULL,
  orchestration_options = NULL,
  performance_profile = NULL,
  webpage_policy = NULL,
  query_templates = NULL,
  verbose = FALSE,
  allow_read_webpages = NULL,
  webpage_relevance_mode = NULL,
  webpage_embedding_provider = NULL,
  webpage_embedding_model = NULL,
  use_plan_mode = FALSE,
  recursion_limit = NULL
)
```

## Arguments

|                        |                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                 |
|------------------------|-----------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| prompt                 | The task prompt or question for the agent to research                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                           |
| output_format          | Expected output format. One of: <ul style="list-style-type: none"> <li>• "text": Returns response text (default)</li> <li>• "json": Parse response as JSON</li> <li>• "raw": Include the raw Python response object for low-level debugging in addition to the normal trace fields</li> <li>• Character vector: Extract specific fields from response</li> </ul>                                                                                                                                                                                                                                                                                                                                                                                                                                                |
| temporal               | Named list or <code>asa_temporal</code> object for temporal filtering: <ul style="list-style-type: none"> <li>• time_filter: DuckDuckGo time filter - "d" (day), "w" (week), "m" (month), "y" (year)</li> <li>• after: ISO 8601 date (e.g., "2020-01-01") - hint for results after this date (added to prompt context)</li> <li>• before: ISO 8601 date (e.g., "2024-01-01") - hint for results before this date (added to prompt context)</li> <li>• strictness: "best_effort" (default) or "strict"</li> <li>• use_wayback: If TRUE, expose Wayback Machine search for historical source checks. Disabled by default.</li> </ul>                                                                                                                                                                              |
| config                 | An <code>asa_config</code> object for unified configuration, or NULL to use defaults                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                            |
| agent                  | An <code>asa_agent</code> object from <code>initialize_agent</code> , or NULL to use the currently initialized agent                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                            |
| expected_fields        | Optional character vector of field names expected in JSON output. When provided, validates that all fields are present and non-null. The result will include a <code>parsing_status</code> field with validation details.                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                       |
| expected_schema        | Optional JSON schema tree used for best-effort repair when the agent output is missing required keys (especially when <code>recursion_limit</code> is reached). This should be a nested list describing the required JSON object/array shape, following the conventions used in <code>asa/tests/testthat/test-langgraph-remainingsteps.R</code> . When provided, this bypasses prompt-based schema inference. Structured terminal payloads automatically include sibling metadata keys <code>&lt;field&gt;_source</code> (URL or NULL) and <code>&lt;field&gt;_confidence</code> (numeric score in <code>[0,1]</code> or NULL) for top-level schema fields, including unresolved fields where sibling values are emitted as NULL, even when those sibling keys are not listed in <code>expected_schema</code> . |
| thread_id              | Optional stable identifier for memory folding sessions. When provided, the same thread ID is reused so folded summaries persist across invocations. Defaults to NULL (new thread each call).                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                    |
| field_status           | Optional per-field extraction ledger seed passed into the LangGraph state. Useful for resuming partially resolved schemas.                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                      |
| budget_state           | Optional tool budget state seed passed into the LangGraph state (e.g., to resume prior progress).                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                               |
| search_budget_limit    | Optional integer maximum number of Search tool calls for this run.                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                              |
| unknown_after_searches | Optional integer threshold after which unresolved fields may be marked as unknown.                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                              |



|                                 |                                                                                                                                                                                                                                                                     |
|---------------------------------|---------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| finalize_on_all_fields_resolved | Optional logical flag. When TRUE, the agent finalizes once all required fields are resolved.                                                                                                                                                                        |
| field_rules                     | Optional named list/dict of field-level extraction rules passed directly to the backend state. Use task-agnostic rule objects; NULL leaves backend defaults unchanged.                                                                                              |
| source_policy                   | Optional named list/dict controlling source selection and evidence-policy behavior in the backend state.                                                                                                                                                            |
| retry_policy                    | Optional named list/dict controlling backend retry behavior for follow-up planning and field resolution.                                                                                                                                                            |
| finalization_policy             | Optional named list/dict controlling terminal finalization behavior in the backend state.                                                                                                                                                                           |
| orchestration_options           | Optional named list/dict of generic orchestration controls (component enable/mode and thresholds). This is passed directly to the backend state and must remain task-agnostic.                                                                                      |
| performance_profile             | Optional orchestration profile. One of: "latency", "balanced", or "quality".                                                                                                                                                                                        |
| webpage_policy                  | Optional structured OpenWebpage policy list with keys: max_open_calls, host_cooldown_seconds, blocked_host_ttl_seconds, open_only_if_score_ge, and parallel_open_limit.                                                                                             |
| query_templates                 | Optional named list/dict of backend query template overrides used when orchestration components generate follow-up searches.                                                                                                                                        |
| verbose                         | Print progress messages (default: FALSE)                                                                                                                                                                                                                            |
| allow_read_webpages             | If TRUE, allows the agent to open and read full webpages (HTML/text) via the OpenWebpage tool. Disabled by default.                                                                                                                                                 |
| webpage_relevance_mode          | Relevance selection for opened webpages. One of: "auto" (default), "lexical", "embeddings". When "embeddings" or "auto" with an available provider, the tool uses vector similarity to pick the most relevant excerpts; otherwise it falls back to lexical overlap. |
| webpage_embedding_provider      | Embedding provider to use for relevance. One of: "auto" (default), "openai", "sentence_transformers".                                                                                                                                                               |
| webpage_embedding_model         | Embedding model identifier. For OpenAI, defaults to "text-embedding-3-small". For sentence-transformers, use a local model name (e.g., "all-MiniLM-L6-v2").                                                                                                         |
| use_plan_mode                   | Logical flag. When TRUE, the agent creates and follows a structured execution plan, updating step status during the run.                                                                                                                                            |
| recursion_limit                 | Optional maximum number of agent steps. Precedence is: per-call value (this argument), then configured default from initialize_agent() / asa_config(), then mode-specific fallback (memory folding: 100; standard agent: 20).                                       |

## Details

This function provides the primary interface for running research tasks. For simple text responses, use `output_format = "text"`. For structured outputs, use `output_format = "json"` or specify

field names to extract. Full trace information is returned via `raw_output` (and `trace_json` when available) for every invocation. For low-level debugging with full underlying response access, use `output_format = "raw"`.

When temporal filtering is specified, the search tool's time filter is temporarily set for this task and restored afterward. Date hints (after/before) are appended to the prompt to guide the agent's search behavior.

## Value

An `asa_result` object with:

- `prompt`: The original prompt
- `message`: The agent's response text
- `parsed`: Parsed output (list for JSON/field extraction, NULL for text/raw)
- `raw_output`: Full agent trace (always included)
- `trace_json`: Structured trace JSON (when available)
- `elapsed_time`: Execution time in minutes
- `status`: "success" or "error"
- `search_tier`: Which search tier was used ("primp", "selenium", etc.)
- `parsing_status`: Validation result (if expected\_fields provided)
- `execution`: Operational metadata (thread\_id, stop\_reason, status\_code, tool budget counters, fold\_count, completion\_gate, verification\_status) plus diagnostics counters from runtime quality guards, candidate and retrieval telemetry, finalization status, artifact status, and orchestration policy metadata.
- `action_ascii`: High-level ASCII action map derived from the trace (also available at `execution$action_ascii`)
- `action_steps`: Parsed high-level action steps (also available at `execution$action_steps`)
- `action_overall`: High-level action summary lines (also available at `execution$action_overall`)
- `langgraph_step_timings`: Per-node LangGraph timings in minutes (also available at `execution$langgraph_step_timings`)
- `fold_stats`: Memory folding diagnostics list
- `raw_response`: Full underlying Python response object (for `output_format = "raw"`)

## See Also

[initialize\\_agent](#), [run\\_task\\_batch](#), [asa\\_config](#), [temporal\\_options](#)

## Examples

```
## Not run:
# Initialize agent first
agent <- initialize_agent(backend = "openai", model = "gpt-4.1-mini")

# Simple text query
result <- run_task(
  prompt = "What is the capital of France?",
  output_format = "text",
  agent = agent
)
print(result$message)

# JSON structured output
```

```

result <- run_task(
  prompt = "Find information about Albert Einstein and return JSON with
            fields: birth_year, death_year, nationality, field_of_study",
  output_format = "json",
  agent = agent
)
print(result$parsed)

# Full traces are always included in asa_result
result <- run_task(
  prompt = "Search for information",
  agent = agent
)
cat(result$raw_output)

# Use output_format = "raw" only when you also need raw_response
result <- run_task(
  prompt = "Search for information",
  output_format = "raw",
  agent = agent
)
str(result$raw_response) # Inspect full underlying Python response

# With temporal filtering (past year only)
result <- run_task(
  prompt = "Find recent AI research breakthroughs",
  temporal = temporal_options(time_filter = "y"),
  agent = agent
)

# With date range hint
result <- run_task(
  prompt = "Find tech companies founded recently",
  temporal = list(
    time_filter = "y",
    after = "2020-01-01",
    before = "2024-01-01"
  ),
  agent = agent
)

# Opt in to Wayback Machine for historical source checks
result <- run_task(
  prompt = "Check what example.com published before 2018",
  temporal = temporal_options(before = "2018-01-01", use_wayback = TRUE),
  agent = agent
)

# Using asa_config for unified configuration
config <- asa_config(
  backend = "openai",
  model = "gpt-4.1-mini",
  temporal = temporal_options(time_filter = "y")
)
result <- run_task("Find recent AI research breakthroughs", config = config)

## End(Not run)

```

run\_task\_batch

*Run Multiple Tasks in Batch***Description**

Executes multiple research tasks, optionally in parallel. Includes a circuit breaker that monitors error rates and pauses execution if errors spike, preventing cascading failures.

**Usage**

```
run_task_batch(
  prompts,
  output_format = "text",
  temporal = NULL,
  config = NULL,
  agent = NULL,
  parallel = FALSE,
  workers = NULL,
  progress = TRUE,
  circuit_breaker = TRUE,
  abort_on_trip = FALSE,
  field_status = NULL,
  budget_state = NULL,
  search_budget_limit = NULL,
  unknown_after_searches = NULL,
  finalize_on_all_fields_resolved = NULL,
  orchestration_options = NULL,
  performance_profile = NULL,
  webpage_policy = NULL,
  allow_read_webpages = NULL,
  webpage_relevance_mode = NULL,
  webpage_embedding_provider = NULL,
  webpage_embedding_model = NULL,
  recursion_limit = NULL
)
```

**Arguments**

|               |                                                                                                                                                                                                                                                                                  |
|---------------|----------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| prompts       | Character vector of task prompts, or a data frame with a 'prompt' column                                                                                                                                                                                                         |
| output_format | Expected output format (applies to all tasks)                                                                                                                                                                                                                                    |
| temporal      | Named list for temporal filtering (applies to all tasks). See <a href="#">run_task</a> for details.                                                                                                                                                                              |
| config        | Optional asa_config object to apply across the batch. When provided, config\$temporal is used if temporal is NULL, and config\$workers is used if workers is NULL. In parallel mode, the config is sent to worker processes so each worker can initialize its own agent session. |
| agent         | An asa_agent object                                                                                                                                                                                                                                                              |
| parallel      | Use parallel processing                                                                                                                                                                                                                                                          |

|                                 |                                                                                                                                                 |
|---------------------------------|-------------------------------------------------------------------------------------------------------------------------------------------------|
| workers                         | Number of parallel workers                                                                                                                      |
| progress                        | Show progress messages                                                                                                                          |
| circuit_breaker                 | Enable circuit breaker for error rate monitoring. When enabled, tracks recent error rates and pauses if threshold exceeded. Default TRUE.       |
| abort_on_trip                   | If TRUE, abort the batch when circuit breaker trips. If FALSE (default), wait for cooldown and continue.                                        |
| field_status                    | Optional named list tracking per-field resolution status. Passed through to <a href="#">run_task</a> .                                          |
| budget_state                    | Optional list tracking search budget consumption across batch items. Passed through to <a href="#">run_task</a> .                               |
| search_budget_limit             | Optional integer maximum number of searches per task. Passed through to <a href="#">run_task</a> .                                              |
| unknown_after_searches          | Optional integer threshold: if a field remains unknown after this many searches, mark it resolved. Passed through to <a href="#">run_task</a> . |
| finalize_on_all_fields_resolved | Optional logical; if TRUE, finalize immediately when all schema fields are resolved. Passed through to <a href="#">run_task</a> .               |
| orchestration_options           | Optional named list/dict of generic orchestration options passed through to <a href="#">run_task</a> .                                          |
| performance_profile             | Optional orchestration profile passed through to <a href="#">run_task</a> .                                                                     |
| webpage_policy                  | Optional structured OpenWebpage policy passed through to <a href="#">run_task</a> .                                                             |
| allow_read_webpages             | If TRUE, allows the agent to open and read full webpages (HTML/text) via the OpenWebpage tool. Disabled by default.                             |
| webpage_relevance_mode          | Relevance selection for opened webpages. One of: "auto", "lexical", "embeddings". See <a href="#">run_task</a> .                                |
| webpage_embedding_provider      | Embedding provider for relevance. See <a href="#">run_task</a> .                                                                                |
| webpage_embedding_model         | Embedding model identifier. See <a href="#">run_task</a> .                                                                                      |
| recursion_limit                 | Optional maximum number of agent steps applied to each task. See <a href="#">run_task</a> .                                                     |

**Value**

A list of `asa_result` objects, or if prompts was a data frame, the data frame with result columns added (including per-row operational metadata and full per-row `asa_result` objects via list column `asa_result` and attribute `asa_results`). If circuit breaker aborts, includes attribute `"circuit_breaker_aborted" = TRUE`.

**See Also**

[run\\_task](#), [configure\\_temporal](#)

## Examples

```
## Not run:
prompts <- c(
  "What is the population of Tokyo?",
  "What is the population of New York?",
  "What is the population of London?"
)
results <- run_task_batch(prompts, agent = agent)

# With temporal filtering for all tasks
results <- run_task_batch(
  prompts,
  temporal = list(time_filter = "y"),
  agent = agent
)

# Opt in to Wayback support for all batch tasks
results <- run_task_batch(
  prompts,
  temporal = temporal_options(before = "2018-01-01", use_wayback = TRUE),
  agent = agent
)

# Disable circuit breaker
results <- run_task_batch(prompts, agent = agent, circuit_breaker = FALSE)

# Abort on circuit breaker trip
results <- run_task_batch(prompts, agent = agent, abort_on_trip = TRUE)

## End(Not run)
```

---

search\_options

Create Search Options

---

## Description

Creates search configuration for controlling DuckDuckGo search behavior, including rate limiting, retry policies, and result limits. These options are used by the 4-tier search fallback system.

## Usage

```
search_options(
  max_results = NULL,
  timeout = NULL,
  max_retries = NULL,
  retry_delay = NULL,
  backoff_multiplier = NULL,
  inter_search_delay = NULL,
  humanize_timing = NULL,
  jitter_factor = NULL,
  allow_direct_fallback = NULL,
  selenium_browser_preference = NULL,
```

```

    stability_profile = NULL,
    performance_profile = NULL,
    webpage_policy = NULL,
    auto_openwebpage_policy = NULL,
    langgraph_node_retries = NULL,
    langgraph_cache_enabled = NULL,
    finalize_when_all_unresolved_exhausted = NULL,
    field_attempt_budget_mode = NULL,
    allow_read_webpages = NULL,
    webpage_relevance_mode = NULL,
    webpage_heuristic_profile = NULL,
    webpage_embedding_provider = NULL,
    webpage_embedding_model = NULL,
    webpage_timeout = NULL,
    webpage_max_bytes = NULL,
    webpage_max_chars = NULL,
    webpage_max_chunks = NULL,
    webpage_chunk_chars = NULL,
    webpage_embedding_api_base = NULL,
    webpage_prefilter_k = NULL,
    webpage_use_mmr = NULL,
    webpage_mmr_lambda = NULL,
    webpage_cache_enabled = NULL,
    webpage_cache_max_entries = NULL,
    webpage_cache_max_text_chars = NULL,
    webpage_blocked_cache_ttl_sec = NULL,
    webpage_blocked_cache_max_entries = NULL,
    webpage_blocked_probe_bytes = NULL,
    webpage_blocked_detect_on_200 = NULL,
    webpage_blocked_body_scan_bytes = NULL,
    webpage_pdf_enabled = NULL,
    webpage_pdf_timeout = NULL,
    webpage_pdf_max_bytes = NULL,
    webpage_pdf_max_pages = NULL,
    webpage_pdf_max_text_chars = NULL,
    webpage_user_agent = NULL,
    wikidata_template_path = NULL,
    wiki_top_k_results = NULL,
    wiki_doc_content_chars_max = NULL,
    search_doc_content_chars_max = NULL
)

```

### Arguments

|             |                                                                                                                                          |
|-------------|------------------------------------------------------------------------------------------------------------------------------------------|
| max_results | Maximum number of search results to return per query. Higher values provide more context but increase latency. Default: 10.              |
| timeout     | Timeout in seconds for individual search requests. Applies to each tier attempt separately. Default: 30.                                 |
| max_retries | Maximum number of retry attempts when a search tier fails. After exhausting retries, the system falls back to the next tier. Default: 3. |
| retry_delay | Initial delay in seconds before the first retry. Subsequent retries use exponential backoff. Default: 2.                                 |

|                                        |                                                                                                                                                                                       |
|----------------------------------------|---------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| backoff_multiplier                     | Multiplier for exponential backoff between retries. E.g., with retry_delay=2 and multiplier=1.5, delays are 2s, 3s, 4.5s. Default: 1.5.                                               |
| inter_search_delay                     | Minimum delay in seconds between consecutive searches. Helps avoid rate limiting from search providers. Default: 1.5.                                                                 |
| humanize_timing                        | Whether to humanize search timing with non-uniform delays. Default uses ASA_HUMANIZE_TIMING.                                                                                          |
| jitter_factor                          | Random jitter factor applied to humanized timing and retry spacing. Default uses ASA_JITTER_FACTOR.                                                                                   |
| allow_direct_fallback                  | Whether to allow direct fallback requests in the search transport. Default: FALSE.                                                                                                    |
| selenium_browser_preference            | Selenium browser engine preference order. One of "firefox_first" (default) or "chrome_first".                                                                                         |
| stability_profile                      | High-level search stability profile. One of "deterministic", "balanced", or "stealth_first" (default). This controls default anti-detection and orchestration behavior.               |
| performance_profile                    | Orchestration profile optimized for one of: "latency", "balanced", or "quality".                                                                                                      |
| webpage_policy                         | Structured OpenWebpage policy list with keys: max_open_calls, host_cooldown_seconds, blocked_host_ttl_seconds, open_only_if_score_ge, and parallel_open_limit.                        |
| auto_openwebpage_policy                | Automatic OpenWebpage follow-up policy. Use "auto" for profile-driven automatic opening or "off" to disable it. Defaults to "off" for the deterministic profile and "auto" otherwise. |
| langgraph_node_retries                 | Whether to enable LangGraph node-level retries in agent graph construction. Default: TRUE.                                                                                            |
| langgraph_cache_enabled                | Whether to enable LangGraph node caching. Default: FALSE.                                                                                                                             |
| finalize_when_all_unresolved_exhausted | Optional override for orchestration finalizer behavior. When TRUE, finalize once all unresolved fields have exhausted their search-attempt budget.                                    |
| field_attempt_budget_mode              | Optional field-attempt budget mode passed to orchestration field_resolver. One of "strict_cap" or "soft_cap".                                                                         |
| allow_read_webpages                    | If TRUE, allows the agent to open and read full webpages (HTML/text) via the OpenWebpage tool. Disabled by default.                                                                   |
| webpage_relevance_mode                 | Relevance selection for opened webpages. One of: "auto", "lexical", "embeddings".                                                                                                     |
| webpage_heuristic_profile              | Heuristic profile used when annotating links from opened webpages. Must be "generic" (task-agnostic default).                                                                         |
| webpage_embedding_provider             | Embedding provider for relevance. One of: "auto", "openai", "azure-openai", "sentence_transformers".                                                                                  |



|                                   |                                                                                                                                                                          |
|-----------------------------------|--------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| webpage_embedding_model           | Embedding model identifier for relevance.                                                                                                                                |
| webpage_timeout                   | Timeout in seconds for OpenWebpage fetches and (when used) relevance embeddings.                                                                                         |
| webpage_max_bytes                 | Maximum bytes to download per page (hard cap).                                                                                                                           |
| webpage_max_chars                 | Maximum characters returned from an opened page (hard cap on tool output).                                                                                               |
| webpage_max_chunks                | Maximum number of relevant excerpts to include from a page.                                                                                                              |
| webpage_chunk_chars               | Approximate size (in characters) of each excerpt.                                                                                                                        |
| webpage_embedding_api_base        | Optional OpenAI-compatible base URL for embeddings. If NULL, uses provider defaults. For Azure embeddings, set AZURE_OPENAI_EMBEDDING_ENDPOINT or AZURE_OPENAI_ENDPOINT. |
| webpage_prefilter_k               | Optional lexical prefilter size before embedding relevance selection.                                                                                                    |
| webpage_use_mmr                   | Whether to apply maximal marginal relevance (MMR) for more diverse excerpts.                                                                                             |
| webpage_mmr_lambda                | MMR tradeoff between relevance (1.0) and diversity (0.0).                                                                                                                |
| webpage_cache_enabled             | Enable per-run caching of opened pages.                                                                                                                                  |
| webpage_cache_max_entries         | Max cached page entries per run.                                                                                                                                         |
| webpage_cache_max_text_chars      | Max characters of extracted page text to store per cached page (separate from webpage_max_chars, which caps tool output).                                                |
| webpage_blocked_cache_ttl_sec     | TTL in seconds for cached blocked URL responses (HTTP 403/429 and anti-bot pages).                                                                                       |
| webpage_blocked_cache_max_entries | Max blocked-URL cache entries per run.                                                                                                                                   |
| webpage_blocked_probe_bytes       | Max bytes to inspect when classifying a response as blocked.                                                                                                             |
| webpage_blocked_detect_on_200     | If TRUE, classify anti-bot interstitials served with HTTP 200 as blocked responses.                                                                                      |
| webpage_blocked_body_scan_bytes   | Max response bytes scanned on HTTP 200 pages for anti-bot marker detection.                                                                                              |
| webpage_pdf_enabled               | If TRUE, OpenWebpage attempts PDF extraction via pdftotext.                                                                                                              |
| webpage_pdf_timeout               | Timeout in seconds for pdftotext.                                                                                                                                        |
| webpage_pdf_max_bytes             | Max PDF bytes to download before extraction.                                                                                                                             |

|                              |                                                                                                                                                                                               |
|------------------------------|-----------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| webpage_pdf_max_pages        | Max pages passed to pdftotext.                                                                                                                                                                |
| webpage_pdf_max_text_chars   | Max extracted PDF characters kept before relevance selection/output truncation.                                                                                                               |
| webpage_user_agent           | User-Agent string used for webpage fetches.                                                                                                                                                   |
| wikidata_template_path       | Optional path to a Wikidata template JSON file. When provided, Wikidata known-entity catalogs are loaded from this file. If NULL (default), no task-specific template catalog is auto-loaded. |
| wiki_top_k_results           | Number of Wikipedia search results to fetch per query (default: 5). Higher values may provide more coverage at the cost of latency/noise.                                                     |
| wiki_doc_content_chars_max   | Maximum number of characters to include from each Wikipedia page/summary in the Wikipedia tool output (default: 1000).                                                                        |
| search_doc_content_chars_max | Maximum number of characters to include from DuckDuckGo Search tool outputs (default: 500). This caps the text returned to the LLM per Search tool invocation.                                |

## Details

The search system uses a 4-tier fallback architecture:

1. **PRIMP**: HTTP/2 with browser TLS fingerprint
2. **Selenium**: Headless browser for JS-rendered content
3. **DDGS**: Standard ddgs Python library
4. **Requests**: Raw POST to DuckDuckGo HTML endpoint

The retry/backoff settings apply within each tier. If all retries are exhausted, the system automatically falls back to the next tier. Selenium browser initialization defaults to "firefox\_first" and falls back to Chrome engines when Firefox is unavailable.

## Value

An object of class `asa_search`

## See Also

[asa\\_config](#), [configure\\_search](#)

## Examples

```
## Not run:
# Default settings
search <- search_options()

# More aggressive settings for faster searches
search <- search_options(
  max_results = 5,
  timeout = 10,
  max_retries = 2
```

```

)

# Conservative settings for rate-limited environments
search <- search_options(
  inter_search_delay = 2.0,
  max_retries = 5,
  backoff_multiplier = 2.0
)

# Use with asa_config
config <- asa_config(
  backend = "openai",
  search = search_options(max_results = 15)
)

## End(Not run)

```

summary.asa\_agent

*Summary Method for asa\_agent Objects***Description**

Summary Method for asa\_agent Objects

**Usage**

```
## S3 method for class 'asa_agent'
summary(object, ...)
```

**Arguments**

|        |                                |
|--------|--------------------------------|
| object | An asa_agent object            |
| ...    | Additional arguments (ignored) |

**Value**

Invisibly returns a summary list

summary.asa\_audit\_result

*Summary Method for asa\_audit\_result Objects***Description**

Summary Method for asa\_audit\_result Objects

**Usage**

```
## S3 method for class 'asa_audit_result'
summary(object, ...)
```

**Arguments**

|        |                                |
|--------|--------------------------------|
| object | An asa_audit_result object     |
| ...    | Additional arguments (ignored) |

**Value**

Invisibly returns a summary list

---

summary.asa\_benchmark    *Summary Method for asa\_benchmark Objects*

---

**Description**

Summary Method for asa\_benchmark Objects

**Usage**

```
## S3 method for class 'asa_benchmark'  
summary(object, ...)
```

**Arguments**

|        |                                |
|--------|--------------------------------|
| object | An asa_benchmark object        |
| ...    | Additional arguments (ignored) |

**Value**

Invisibly returns a benchmark summary list

---

summary.asa\_enumerate\_result  
                                  *Summary Method for asa\_enumerate\_result Objects*

---

**Description**

Summary Method for asa\_enumerate\_result Objects

**Usage**

```
## S3 method for class 'asa_enumerate_result'  
summary(object, ...)
```

**Arguments**

|        |                                |
|--------|--------------------------------|
| object | An asa_enumerate_result object |
| ...    | Additional arguments (ignored) |

**Value**

Invisibly returns a summary list

---

`summary.asa_evaluation_result`*Summary Method for asa\_evaluation\_result Objects*

---

**Description**

Summary Method for asa\_evaluation\_result Objects

**Usage**

```
## S3 method for class 'asa_evaluation_result'  
summary(object, ...)
```

**Arguments**

|                     |                                 |
|---------------------|---------------------------------|
| <code>object</code> | An asa_evaluation_result object |
| <code>...</code>    | Additional arguments (ignored)  |

**Value**

Invisibly returns an evaluation summary list

---

`summary.asa_response`    *Summary Method for asa\_response Objects*

---

**Description**

Summary Method for asa\_response Objects

**Usage**

```
## S3 method for class 'asa_response'  
summary(object, show_trace = FALSE, ...)
```

**Arguments**

|                         |                                |
|-------------------------|--------------------------------|
| <code>object</code>     | An asa_response object         |
| <code>show_trace</code> | Include full trace in output   |
| <code>...</code>        | Additional arguments (ignored) |

**Value**

Invisibly returns a summary list

---

|                    |                                              |
|--------------------|----------------------------------------------|
| summary.asa_result | <i>Summary Method for asa_result Objects</i> |
|--------------------|----------------------------------------------|

---

**Description**

Summary Method for asa\_result Objects

**Usage**

```
## S3 method for class 'asa_result'
summary(object, ...)
```

**Arguments**

|        |                                |
|--------|--------------------------------|
| object | An asa_result object           |
| ...    | Additional arguments (ignored) |

**Value**

Invisibly returns a summary list

---

|                  |                                          |
|------------------|------------------------------------------|
| temporal_options | <i>Create Temporal Filtering Options</i> |
|------------------|------------------------------------------|

---

**Description**

Creates a temporal filtering configuration for constraining search results by date. Supports DuckDuckGo time filters, date ranges, and strict verification modes.

**Usage**

```
temporal_options(
  time_filter = NULL,
  after = NULL,
  before = NULL,
  strictness = "best_effort",
  use_wayback = FALSE
)
```

**Arguments**

|             |                                                                                               |
|-------------|-----------------------------------------------------------------------------------------------|
| time_filter | DuckDuckGo time filter: "d" (day), "w" (week), "m" (month), "y" (year), or NULL for no filter |
| after       | ISO 8601 date string (e.g., "2020-01-01") - results after this date                           |
| before      | ISO 8601 date string (e.g., "2024-01-01") - results before this date                          |
| strictness  | Verification level: "best_effort" (default) or "strict"                                       |
| use_wayback | Use Wayback Machine for strict pre-date guarantees                                            |

**Details**

Temporal filtering can operate at different levels:

- **time\_filter**: DuckDuckGo native filter (fast, approximate)
- **after/before**: Date hints appended to prompts
- **strict**: Post-hoc verification of result dates
- **use\_wayback**: Uses Internet Archive for guaranteed historical data

**Value**

An object of class `asa_temporal`

**See Also**

[asa\\_config](#), [run\\_task](#)

**Examples**

```
## Not run:
# Past year only
temporal <- temporal_options(time_filter = "y")

# Specific date range
temporal <- temporal_options(
  after = "2020-01-01",
  before = "2024-01-01"
)

# Strict historical verification
temporal <- temporal_options(
  before = "2015-01-01",
  strictness = "strict",
  use_wayback = TRUE
)

## End(Not run)
```

---

tor\_options

*Tor Options*


---

**Description**

Configure shared Tor exit tracking for healthier circuit rotation.

**Usage**

```
tor_options(
  registry_path = NULL,
  dirty_tor_exists = ASA_TOR_REGISTRY_ENABLED,
  bad_ttl = ASA_TOR_BAD_TTL,
  good_ttl = ASA_TOR_GOOD_TTL,
```

```

    overuse_threshold = ASA_TOR_OVERUSE_THRESHOLD,
    overuse_decay = ASA_TOR_OVERUSE_DECAY,
    max_rotation_attempts = ASA_TOR_MAX_ROTATION_ATTEMPTS,
    ip_cache_ttl = ASA_TOR_IP_CACHE_TTL
)

```

### Arguments

|                       |                                                                     |
|-----------------------|---------------------------------------------------------------------|
| registry_path         | Path to the shared SQLite registry file (default: user cache).      |
| dirty_tor_exists      | Enable the registry (tracks good/bad/overused exits).               |
| bad_ttl               | Seconds to keep a bad/tainted exit before reuse (default: 3600).    |
| good_ttl              | Seconds to treat an exit as good before refreshing (default: 1800). |
| overuse_threshold     | Max recent uses before a good exit is considered overloaded.        |
| overuse_decay         | Window (seconds) for overuse counting before decaying.              |
| max_rotation_attempts | Max attempts to find a clean exit before giving up.                 |
| ip_cache_ttl          | Seconds to cache exit IP lookups.                                   |

### Value

An object of class `asa_tor`

---

|                                |                                          |
|--------------------------------|------------------------------------------|
| write_csv.asa_enumerate_result | <i>Write asa_enumerate_result to CSV</i> |
|--------------------------------|------------------------------------------|

---

### Description

Write `asa_enumerate_result` to CSV

### Usage

```
write_csv.asa_enumerate_result(x, file, include_provenance = FALSE, ...)
```

### Arguments

|                    |                                                       |
|--------------------|-------------------------------------------------------|
| x                  | An <code>asa_enumerate_result</code> object           |
| file               | Path to output CSV file                               |
| include_provenance | Include provenance as additional columns              |
| ...                | Additional arguments passed to <code>write.csv</code> |

### Value

Invisibly returns the file path



# Index

`as.data.frame.asa_audit_result`, 3  
`as.data.frame.asa_enumerate_result`, 4  
`as.data.frame.asa_evaluation_result`, 4  
`as.data.frame.asa_result`, 5  
`asa_agent`, 5  
`asa_audit`, 6  
`asa_audit_result`, 8  
`asa_config`, 9, 50, 58, 63  
`asa_enumerate`, 11, 27, 35  
`asa_enumerate_result`, 15  
`asa_evaluation_result`, 16  
`asa_response`, 17  
`asa_result`, 18  
  
`benchmark_from_df`, 19  
`benchmark_from_yaml`, 20  
`build_backend`, 21  
`build_prompt`, 22  
  
`check_backend`, 23  
`configure_search`, 24, 58  
`configure_search_logging`, 25  
`configure_temporal`, 26, 53  
`configure_tor_registry`, 27  
`create_benchmark`, 28  
  
`evaluate_agent`, 29  
`extract_agent_results`, 29  
`extract_search_snippets`, 30  
`extract_search_tiers`, 31  
`extract_urls`, 32  
`extract_wikipedia_content`, 32  
  
`get_agent`, 33  
`get_tor_ip`, 33  
  
`initialize_agent`, 14, 34, 47, 48, 50  
`is_tor_running`, 37  
  
`list_benchmarks`, 38  
`load_benchmark`, 38  
  
`print.asa_agent`, 39  
`print.asa_audit_result`, 39  
`print.asa_benchmark`, 40  
  
`print.asa_config`, 40  
`print.asa_enumerate_result`, 41  
`print.asa_evaluation_result`, 41  
`print.asa_response`, 42  
`print.asa_result`, 42  
`print.asa_search`, 43  
`print.asa_temporal`, 43  
`print.asa_tor`, 44  
`process_outputs`, 44  
  
`reset_agent`, 45  
`rotate_tor_circuit`, 45  
`run_direct_task`, 46  
`run_task`, 10, 14, 19, 27, 35, 37, 47, 52, 53, 63  
`run_task_batch`, 10, 37, 50, 52  
  
`search_options`, 10, 35, 54  
`summary.asa_agent`, 59  
`summary.asa_audit_result`, 59  
`summary.asa_benchmark`, 60  
`summary.asa_enumerate_result`, 60  
`summary.asa_evaluation_result`, 61  
`summary.asa_response`, 61  
`summary.asa_result`, 62  
  
`temporal_options`, 10, 50, 62  
`tor_options`, 36, 63  
  
`write_csv.asa_enumerate_result`, 64